

Advance Network Programming

Processes

Outline

- Process definition
- Process handling in Unix
- Process creation
- Process termination
- Process synchronization

Process Definition

- Definition: A process is an instance of a running program.
 - One of the most fundamental concepts in computer science.
 - Not the same as “program” or “processor”
- A program is a set of instructions and initialized data in a file, usually found on a disk.
- A process is an instance of that program while it is running, along with the state of all the CPU registers and the values of data in memory.
- A single program can correspond to many processes; for example, several users can be running a shell.

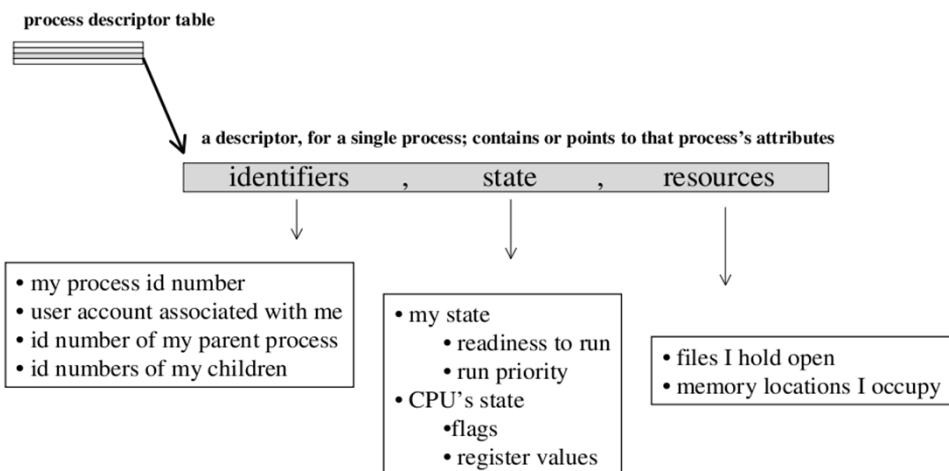
Process Handling

- Constituents of a process
 - Its code
 - Data
 - its own
 - OS’s data used by/for process
 - Various attributes OS needs to manage it

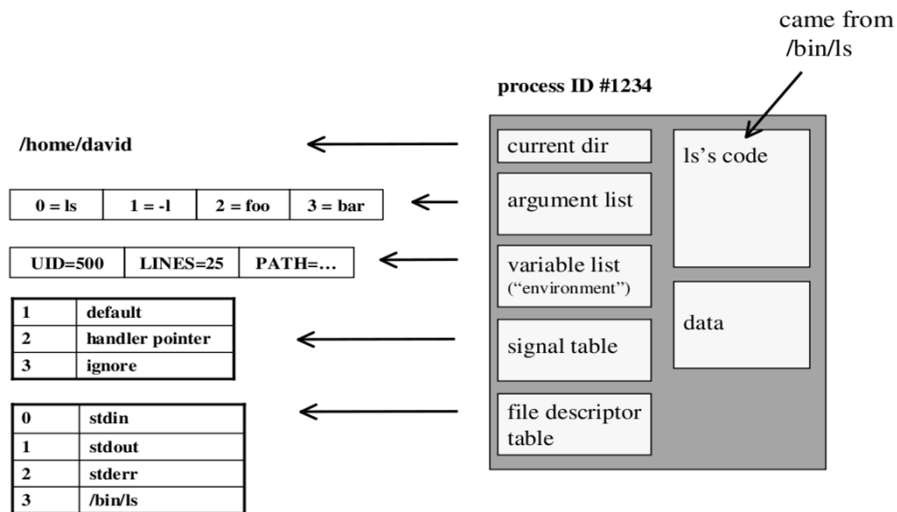
Process Handling

- OS keeps track of all processes
 - Process table/array/list
 - Elements are process descriptors (aka control blocks)
 - Descriptors reference code & data

Process Descriptor



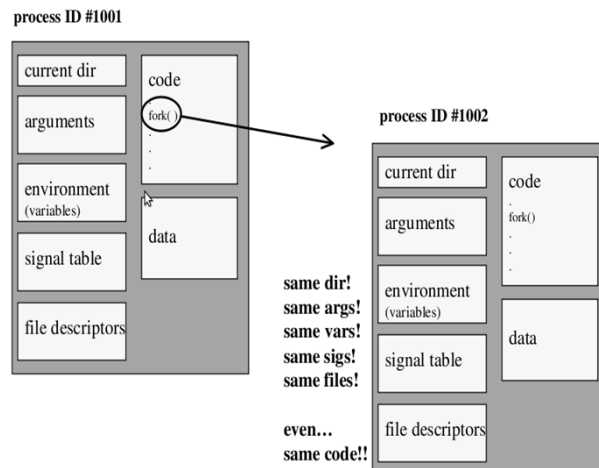
ls -l foo bar



Process Creation in Unix

- `fork()` function
- `exec()` family of functions
- `spawn()`
- `wait()`
- `exit()`

fork() system call



fork() example

```
File Edit View Terminal Tabs Help
cihan@sdf-1:...ogramming/course/2
cihan@sdf-1:/media/My Book/Linux-28.02.2009/networkProgramming/course/2> cat Fork1.c
#include <stdio.h>
#include <unistd.h>

int main ( void ) {

    printf("Message before fork\n");

    fork();

    printf("Message after fork\n");

    return 0;
}

cihan@sdf-1:/media/My Book/Linux-28.02.2009/networkProgramming/course/2> ./Fork1
Message before fork
Message after fork
Message after fork
cihan@sdf-1:/media/My Book/Linux-28.02.2009/networkProgramming/course/2>
```

Double output

Single execution

Process Differentiation

- identical? not what we had in mind!
- more useful if child does different stuff
- can we give it different behavior?

provide different behavior

- in the form of source code
- in the form of an existing binary executable
 - exec() family of functions

fork - how to self-identify?

```
cihan@sdf-1:...ogramming/course/2
File Edit View Terminal Tabs Help
cihan@sdf-1:/media/My Book/Linux-28.02.2009/networkProgramming/course/2> cat Fork2.c
#include <stdio.h>
#include <unistd.h>

int main ( void ) {

    int forkResult;

    printf("process id : %i\n",getpid());
    forkResult = fork();
    printf("process id : %i - result : %d\n",getpid(), forkResult);

    return 0;
}

cihan@sdf-1:/media/My Book/Linux-28.02.2009/networkProgramming/course/2> ./Fork2
process id : 308
process id : 309 - result : 0
process id : 308 - result : 309
cihan@sdf-1:/media/My Book/Linux-28.02.2009/networkProgramming/course/2>
```

If 0, I must be the child copy

If not, I must be the parent copy

By source code

```
cihan@sdf-1:...ogramming/course/2
File Edit View Terminal Tabs Help
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main ( void ) {

    printf("(%) Parent does something...\n", getpid());

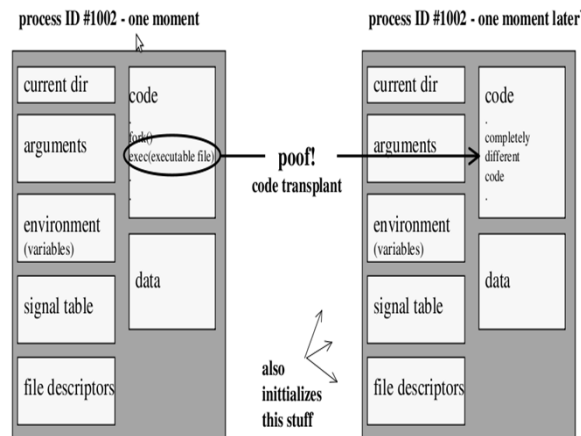
    if(fork()) { // Parent
        printf("(%) Parent do completely different stuff\n",getpid());
    } else { // Child
        printf("(%) Child can do some stuff\n",getpid());
    }

    exit(0);
}

cihan@sdf-1:...ogramming/course/2
File Edit View Terminal Tabs Help
cihan@sdf-1:/media/My Book__/_Linux-28.02.2009/networkProgramming/course/2> ./Fork3
(6872) Parent does something...
(6873) Child can do some stuff
(6872) Parent do completely different stuff
cihan@sdf-1:/media/My Book__/_Linux-28.02.2009/networkProgramming/course/2>
```

conditional, on whether parent or child

By *exec()* function



exec() example

```
File Edit View Terminal Tabs Help
cihan@sdf-1:...ogramming/course/2

#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main ( void ) {
    printf("Parent does stuff and then calls fork...\n");
    if(fork()) { // Parent
        printf("... parent do something completely different\n");
    } else { // Child
        printf("Child runs an executable...\n");
        exec("bin/ls", "bin/ls", "-l", "/etc/apache2/conf.d/", NULL);
    }
    exit(0);
}
```

```
File Edit View Terminal Tabs Help
cihan@sdf-1:...ogramming/course/2

cihan@sdf-1:/media/My Book___/Linux-28.02.2009/networkProgramming/course/2> ./Exec
Parent does stuff and then calls fork...
Child runs an executable...
total 16
-rw-r--r-- 1 root root 646 2008-12-03 12:14 apache2-manual.conf
-rw-r--r-- 1 root root 225 2008-12-04 09:24 gitweb.conf
-rw-r--r-- 1 root root 709 2008-12-03 12:48 mod_perl.conf
-rw-r--r-- 1 root root 451 2008-12-03 13:07 php5.conf
... parent do something completely different
cihan@sdf-1:/media/My Book___/Linux-28.02.2009/networkProgramming/course/2>
```

ls -l /etc/apache2/conf.d/

exec() Family of Functions

- `int execl(const char *pathname, const char *arg0, ...);`
- `int execv(const char *pathname, char *const argv[]);`
- `int execl(const char *pathname, const char *arg0, ..., 0, char *const envp[]);`
- `int execlp(const char *filename, const char *arg0, ...);`
- `int execvp(const char *filename, char *const argv[]);`
- `int execve(const char *pathname, char *const argv[], char *const envp[]);`

For example...

- Shell is running
- You type “ls” and Enter
- Shell is parent, spawns ls as child

A simple shell example

```
int main(void)
{
    char inputBuffer[MAX_LINE]; /* buffer to hold the command entered */
    int background;             /* equals 1 if a command is followed by '&' */
    char *args[MAX_LINE/2+1]; /* command line (of 80) has max of 40 arguments */

    while (1){                  /* Program terminates normally inside setup */
        background = 0;
        printf("COMMAND->");
        fflush(0);
        setup(inputBuffer, args, &background); /* get next command */

        /* the steps are:
        (1) fork a child process using fork()
        (2) the child process will invoke execvp()
        (3) if background == 0, the parent will wait,
            otherwise returns to the setup() function. */
    }
}

~
~
"shell.c" 87L, 2813C written
```

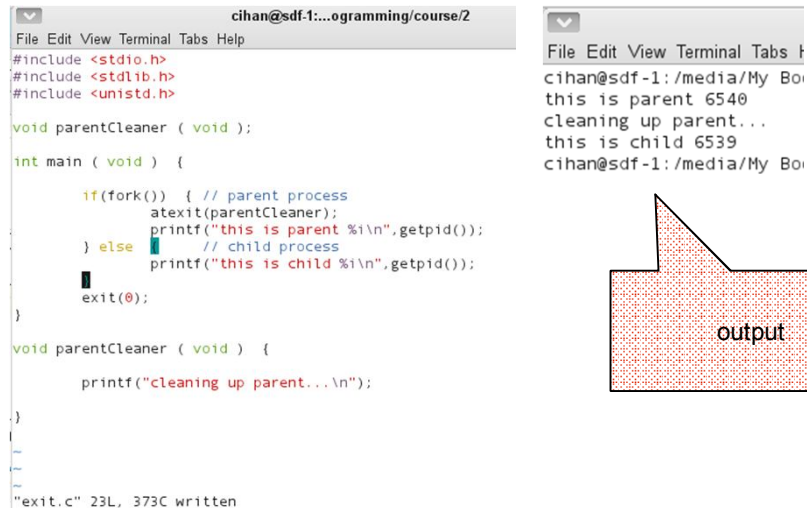
80,0-1

Bot

Process Termination

- void exit (int status);
 - exits a process
 - normally return with status 0
- int atexit (void (*function)(void));
 - registers function to be executed on exit

Process Termination Example



The image shows a code editor window on the left and a terminal window on the right. The code editor displays a C program named 'exit.c' with the following content:

```
File Edit View Terminal Tabs Help
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

void parentCleaner ( void );

int main ( void ) {

    if(fork()) { // parent process
        atexit(parentCleaner);
        printf("this is parent %i\n",getpid());
    } else { // child process
        printf("this is child %i\n",getpid());
    }
    exit(0);
}

void parentCleaner ( void ) {
    printf("cleaning up parent...\n");
}

"exit.c" 23L, 373C written
```

The terminal window on the right shows the output of the program:

```
File Edit View Terminal Tabs Help
cihan@sdf-1: /media/My Bo
this is parent 6540
cleaning up parent...
this is child 6539
cihan@sdf-1: /media/My Bo
```

A red speech bubble with the word "output" points to the terminal window.

Zombie Processes

- When process terminates, still consumes system resources
 - Various tables maintained by OS
- Called a zombie
 - Living corpse, half alive, half dead

Zombie Processes

- Reaping
 - Performed by parent on terminated child
 - Parent is given exit status information
 - Kernel discards process
- What if parent does not reap ?
 - if any parent terminates without reaping a child, then child will be reaped by init process
 - so, only need explicit reaping in long-running processes

Zombie : non-terminating parent

```
File Edit View Terminal Tabs Help
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main ( void ) {

    if(fork(0)) { // Parent
        printf("Running parent, pid : %i\n",getpid());
        while(1);
    } else { // Child
        printf("Terminating child, pid : %i\n", getpid());
        exit(0);
    }

    exit (0);
}
```

```
File Edit View Terminal Tabs Help
cihan@sdf-1: /media/My Book/Linux-28.02.2009/networkProgramming/course/2> ./Zombie.c
Terminating child, pid : 7423
Running parent, pid : 7422
[1] 7422
cihan@sdf-1: /media/My Book/Linux-28.02.2009/networkProgramming/course/2> ps
  PID TTY          TIME CMD
 7197 pts/3    00:00:00 bash
 7422 pts/3    00:00:02 Zombie
 7423 pts/3    00:00:00 Zombie <defunct>
 7427 pts/3    00:00:00 ps
cihan@sdf-1: /media/My Book/Linux-28.02.2009/networkProgramming/course/2>
```

Ps shows chil as defunct

Killing parent allows child
to be reaped by shell

Zombie : non-terminating child

```
File Edit View Terminal Tabs Help
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>

int main ( void ) {

    if(fork()) { // Parent
        printf("Running parent, pid : %i\n", getpid());
        exit(0);
    } else { // Child
        printf("Terminating child, pid : %i\n", getpid());
        while(1);
    }

    exit (0);
}
```

Child process still active, even though parent has terminated

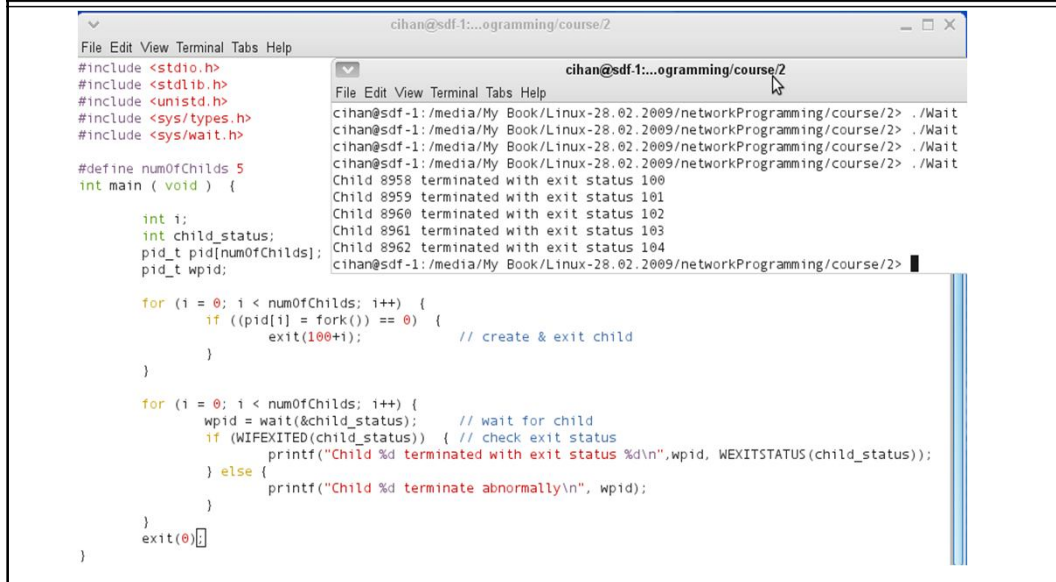
```
File Edit View Terminal Tabs Help
cihan@sdf-1:/media/My Book/Linux-28.02.2009/networkProgramming/course/2> ./Zombie2
Terminating child, pid : 7911
Running parent, pid : 7910
cihan@sdf-1:/media/My Book/Linux-28.02.2009/networkProgramming/course/2> ps
PID TTY TIME CMD
7197 pts/3 00:00:00 bash
7911 pts/3 00:00:01 Zombie2
7918 pts/3 00:00:00 ps
cihan@sdf-1:/media/My Book/Linux-28.02.2009/networkProgramming/course/2>
```

Must be killed explicitly, or Else will keep running indefinitely

Synchronizing with child

- `int wait(int *child_status)`
 - suspends current process until one of its children terminates
 - return value is the pid of the child process that terminated
 - If the child has already terminated, then wait returns its pid immediately
 - If `child_status != NULL`, then the object it points to will be set to a status indicating why the child process terminated

wait() example



```
File Edit View Terminal Tabs Help
cihan@sdf-1:...ogramming/course/2

#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/wait.h>

#define numOfChlds 5
int main ( void ) {
    int i;
    int child_status;
    pid_t pid[numOfChlds];
    pid_t wpid;

    for ( i = 0; i < numOfChlds; i++ ) {
        if ( (pid[i] = fork()) == 0 ) {
            exit(100+i);
        }
    }

    for ( i = 0; i < numOfChlds; i++ ) {
        wpid = wait(&child_status);
        if ( WIFEXITED(child_status) ) { // wait for child
            printf("Child %d terminated with exit status %d\n",wpid, WEXITSTATUS(child_status));
        } else {
            printf("Child %d terminate abnormally\n", wpid);
        }
    }
    exit(0)
}
```

references

- man pages
- <http://www.cs.princeton.edu/courses/archive/fall01/cs217/slides/process.pdf>
- <http://www.cs.cmu.edu/afs/cs.cmu.edu/academic/class/15213-f08/www/lectures/lecture-11.pdf>
- <http://csapp.cs.cmu.edu>
- http://homepage.smc.edu/morgan_david/linux/a12-processes.pdf