


Operating Systems

Interprocess Communication (IPC)

IPC

- Unix
 - pipes, fifos (named pipes)
- System V
 - signals
 - message queues
 - semaphores
 - shared memory
- Posix
 - signals
 - message queues
 - semaphores
 - shared memory

Persistence of IPC Objects



- process-persistent IPC:
 - exists until last process with IPC object closes the object
- kernel-persistent IPC
 - exists until kernel reboots or IPC object is explicitly deleted
- filesystem-persistent IPC
 - exists until IPC object is explicitly deleted

Pipes

- A pipe provides a one-way flow of data
 - example: `who | sort | lpr`
- The difference between a file and a pipe:
 - pipe is a data structure in the kernel.
- A pipe is created by using the pipe system call
 - `int pipe(int* filedes);`
 - Two file descriptors are returned
 - `filedes[0]` is open for reading
 - `filedes[1]` is open for writing
- Typical size is 512 bytes (Minimum limit defined by POSIX)

Pipe Example - II

```

cihan@sdf.t...ogramming/course/3
File Edit View Terminal Tabs Help
#define MAX_LINE 80

void client(int, int);
void server(int, int);

int main (int argc, char *argv[])
{
    int pipe1[2], pipe2[2];
    pid_t childpid;

    pipe(pipe1);
    pipe(pipe2);

    if((childpid=fork())==0) { // Child
        close(pipe1[1]);
        close(pipe2[0]);
        server(pipe1[0], pipe2[1]);
        exit(0);
    }

    close(pipe1[0]);
    close(pipe2[1]);

    client(pipe2[0], pipe1[1]);

    waitpid(childpid, NULL, 0); // wait for child to terminate
    exit(0);
}

void client (int readfd, int writefd) {
    size_t len;
    size_t n;
    char buff[MAX_LINE];

    fgets(buff, MAX_LINE, stdin);
    len = strlen(buff);
    if(buff[len-1]!='\n')
        len--;
    write(writefd, buff, len);

    while((n = read(readfd, buff, MAX_LINE))>0)
        write(STDOUT_FILENO, buff, n);
}

void server(int readfd, int writefd) {
    int fd;
    size_t n;
    char buff[MAX_LINE+1];

    //read path from IPC channel
    if((n = read(readfd, buff, MAX_LINE)) == 0) {
        write(writefd, "EOF while reading path...\n", 26);
        exit(0);
    }
    buff[n] = '\0';
    if((fd=open(buff, O_RDONLY))<0) {
        sprintf(buff+n, sizeof(buff)-n, "can't open, %s\n", strerror(errno));
        n = strlen(buff);
        write(writefd, buff, n);
    } else {
        while((n=read(fd, buff, MAX_LINE))>0)
            write(writefd, buff, n);
        close(fd);
    }
}

```

More pipe functions

- FILE *popen (const char * command, const char *type)
 - Type is r, the calling process reads the standart output of the command,
 - Type is w, the calling process writes to the standart input of the command
 - Return file * if OK, NULL on error
- int pclose (FILE *stream);
 - Closes a standard I/O stream that was created by popen

FIFOs

- Pipes have no names, they can only be used between processes that have a parent process in common.
- FIFO stands for first-in, first-out
- Similar to a pipe, it is a one-way (half duplex) flow of data
- A FIFO has a pathname associated with it, allowing unrelated processes to access a single FIFO
- FIFOs are also called *named pipes*

FIFOs

- `#include <sys/types.h>`
- `#include <sys/stat.h>`
- `int mkfifo(const char *pathname, mode_t mode)`
 - returns 0 if OK, -1 on error

FIFO example

```

File Edit View Terminal Tabs Help
#define MAX_LINE 80
#define FIFO1 "/tmp/fifo.1"
#define FIFO2 "/tmp/fifo.2"

void client(int, int);

int main (int argc, char *argv[]) {
    int readfd, writefd;

    writefd = open(FIFO1, O_WRONLY);
    readfd = open(FIFO2, O_RDONLY);

    client(readfd, writefd);

    close(readfd);
    close(writefd);

    unlink(FIFO1);
    unlink(FIFO2);

    exit(0);
}

File Edit View Terminal Tabs Help
#define MAX_LINE 80
#define FIFO1 "/tmp/fifo.1"
#define FIFO2 "/tmp/fifo.2"
#define FIFO_MODE (S_IRWXU | S_IRWXG | S_IRWXO)

void server(int, int);

int main (int argc, char *argv[]) {
    int readfd, writefd;

    if((mkfifo(FIFO1, FIFO_MODE)<0)&&(errno != EEXIST)) {
        printf("can not open %s\n", FIFO1);
        exit(-1);
    }

    if((mkfifo(FIFO2, FIFO_MODE)<0)&&(errno != EEXIST)) {
        printf("can not open %s\n", FIFO2);
        exit(-1);
    }

    readfd = open(FIFO1, O_RDONLY);
    writefd = open(FIFO2, O_WRONLY);

    server(readfd, writefd);

    exit(0);
}

```

Signals

- Definition
- Signal Types
- Generating Signals
- Responding to a Signal
- POSIX Signal Functions
- Signals & System Calls


Definition

- A signal is an *asynchronous* event which is delivered to a process
- Asynchronous means that the event can occur at any time
 - may be unrelated to the execution of the process
 - e.g. user types `ctrl-C`, or the modem hangs

Common use of Signals

- Ignore a Signal
- Clean up and Terminate
- Dynamic Reconfiguration
- Report Status
- Turn Debugging on/off
- Restore Previous Handler
- Signals & System Calls

Signal Sources



Generating a Signal

- Use the Unix command
 - `$> kill -KILL 4481`
 - Send a SIGKILL signal to pid 4481
 - `ps -l`
 - To make sure process died
- kill function


```
#include <sys/types.h>
#include <signal.h>
int kill(pid_t pid, int sig);
```

Responding to a Signal

- A process can;
 - Ignore/discard the signal (not possible for SIGKILL & SIGSTOP)
 - Execute a signal handler function, and then possibly resume execution or terminate
 - Carry out default action for that signal
- The **choice** is called the process' *signal disposition*

POSIX Signal System

- The POSIX signal system, uses signal sets, to deal with pending signals that might otherwise be missed while a signal is being processed
- The signal set stores collection of signal types
- Sets are used by signal functions to define which signal types are to be processed
- POSIX contains several functions for creating, changing and examining signal sets

POSIX Functions

```
#include <signal.h>

int sigemptyset( sigset_t *set );
int sigfillset( sigset_t *set );
int sigismember( const sigset_t *set,
int signo );
int sigaddset( sigset_t *set, int signo );
int sigdelset( sigset_t *set, int signo );
int sigprocmask ( int how, const sigset_t *set,
sigset_t *oldset);
```

A Critical Code Region

```
• sigset_t newmask, oldmask;
• sigemptyset( &newmask );
• sigaddset( &newmask, SIGINT );
• /* block SIGINT; save old mask */
• sigprocmask( SIG_BLOCK, &newmask, &oldmask );
• /* critical region of code */
• /* reset mask which unblocks SIGINT */
• sigprocmask( SIG_SETMASK, &oldmask, NULL );
```

Example

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <signal.h>

void ouch( int );

int main (void) {
    struct sigaction act;
    act.sa_handler = ouch;
    sigemptyset(&act.sa_mask);
    act.sa_flags = 0;
    sigaction(SIGINT, &act, 0);

    while(1) {
        printf("Hello world\n");
        sleep(1);
    }
    exit 0;
}

void ouch( int sigNo ) {
    printf("received SIGINT...\n");
}
```

- This function will continually capture the ctrl-C (SIGINT) signal.
- Default behavior is **not** restored after signal is caught.
- To terminate the program, must type ctrl-\, the SIGQUIT signal.

Interrupted System Calls

- When a system call is interrupted by a signal, a signal handler is called, returns, and then what ?
- Slow system function calls do not resume, Instead they return an error and errno is assigned EINTR.
- Some UNIXs resume non-slow system functions after the handler has finished.
- Some UNIXs only call the handler after non-slow system function call has finished.

System Calls Inside Handlers

- If a system function is called inside a signal handler then it may interact with an interrupted call to the same function in the main code.
 - e.g. `malloc()`
- Not a problem if the function is reentrant
 - A process can contain multiple calls to these functions at the same time. e.g. `read()`, `write()`
- A function may be non-reentrant for a number of reasons
 - It uses a static data structure
 - It manipulates the heap. e.g. `malloc()`, `free()`
 - It uses standart I/O library. e.g. `printf()`

Message Queues

- Unlike pipes and FIFOs, message queues support messages that have structure.
- Like FIFOs, message queues are persistent objects that must be initially created and eventually deleted when no longer required.
- Message queues are created with a specified maximum message size and maximum number of messages.
- Message queues are created and opened using a special version of the open system call, `mq_open`.

POSIX Message Queue Functions

- `mq_open()`
- `mq_close()`
- `mq_unlink()`
- `mq_send()`
- `mq_receive()`
- `mq_setattr()`
- `mq_getattr()`
- `mq_notify()`

`mq_open(const char *name, int oflag,...)`

- `name`
 - Must start with a slash and contain no other slashes
 - QNX puts these in the `/dev/mqueue` directory
- `oflag`
 - `O_CREAT` – to create a new message queue
 - `O_EXCL` – causes creation to fail if queue exists
 - `O_NONBLOCK` – usual interpretation
- `mode` – usual interpretation
- `&mqattr` – address of structure used during creation

Message Queue Persistence - I

- As noted, a message queue is persistent.
- Unlike a FIFO, however, the contents of a message queue are also persistent.
- It is not necessary for a reader and a writer to have the message queue open at the same time. A writer can open (or create) a queue and write messages to it, then close it and terminate.
- Later a reader can open the queue and read the messages.

Message Queue Persistence - II

- `mkdir /dev/mqueue`
- `mount -t mqueue none /dev/mqueue`
- `ls -la /dev/mqueue`

```
cihan@sdf-1:~/Desktop> ls -la /dev/mqueue/
total 0
drwxrwxrwt 2 root root 80 2009-03-18 20:59 █
drwxr-xr-x 14 root root 4600 2009-03-18 20:52 ..
-rwxr-x--- 1 cihan users 80 2009-03-18 19:40 myqueue123
-rwxr-x--- 1 cihan users 80 2009-03-18 20:59 test
cihan@sdf-1:~/Desktop>
```

Message Queue Example

- `./DropOne -q -p 11`
- `./DropOne -p 101`
- `./DropOne -p 99`

```
File Edit View Terminal Tabs Help
-rwxr-x--- 1 cihan users 80 2009-03-18 21:56 test
cihan@sdf-1:/media/KINGSTON/networkProgramming/course/4> ./TakeOne
Queue "/test":
- stores at most 10 messages
- large at most 8192 bytes each
- currently holds 3 messages
Received message (57 bytes) from 101: Hello from process 14198 (at Wed Mar 18 21:56:26 2009
).
cihan@sdf-1:/media/KINGSTON/networkProgramming/course/4> ./TakeOne
Queue "/test":
- stores at most 10 messages
- large at most 8192 bytes each
- currently holds 2 messages
Received message (57 bytes) from 99: Hello from process 14194 (at Wed Mar 18 21:56:24 2009
).
cihan@sdf-1:/media/KINGSTON/networkProgramming/course/4> ./TakeOne
Queue "/test":
- stores at most 10 messages
- large at most 8192 bytes each
- currently holds 1 messages
Received message (57 bytes) from 11: Hello from process 14187 (at Wed Mar 18 21:56:21 2009
).
cihan@sdf-1:/media/KINGSTON/networkProgramming/course/4> █
```

The effect of fork on a message queue

- Message queue descriptors are not (in general) treated as file descriptors; the unique open, close, and unlink calls should already suggest this.
- Open message queue descriptors are not inherited by child processes created by fork.
- Instead, a child process must explicitly open (using `mq_open`) the message queue itself to obtain a message queue descriptor

Detecting non-empty queues

- `mq_receive` on an empty queue normally causes a process to block, and this may not be desirable.
- Of course, `O_NONBLOCK` could be applied to the queue to prevent this behavior, but in that case the `mq_receive` call will return -1, and our only recourse is to try `mq_receive` again later.
- With the `mq_notify` call we can associate a single process with a message queue so that it (the process) will be notified when the message queue changes state from empty to non-empty

`mq_notify(mqd_t mqdes, const struct sigevent *notification)`

- `queuefd` – as usual, to identify the message queue
- `sigev` – a struct `sigevent` object that identifies the signal to be sent to the process to notify it of the queue state change.
- Once notification has been sent, the notification mechanism is removed. That is, to be notified of the next state change (from empty to non-empty), the notification must be reasserted.

Changing the process to be notified

- Only one process can be registered (at a time) to receive notification when a message is added to a previously-empty queue.
- If you wish to change the process that is to be notified, you must remove the notification from the process which is currently associated (call `mq_notify` with `NULL` for the `sigev` argument), and then associate the notification with a different process.

`mq_getattr(queuefd,&mqstat)`

- This function retrieves the set of attributes for a message queue to the struct `mq_attr` object named `mqstat`.
- Recall that the `mq_flags` member of the attributes is not significant during `mq_open`, but it can be set later (using `mq_setattr`).

`mq_setattr(queuefd,&mqstat,&old)`

- This function is used to
 - Set (or clear) to `O_NONBLOCK` flag in the `mqattr` structure for the identified message queue
 - Retrieve (if `old` is not `NULL`) the previously existing message queue attributes
- Making changes to any other members of the `mqattr` structure is ineffective.

Timed send and receive

- Two additional functions, `mq_timedsend` and `mq_timedreceive`, are like `mq_send` and `mq_receive` except they have an additional argument, a pointer to a `struct timespec`.
- This provides the absolute time at which the send or receive will be aborted if it cannot be completed (because the queue is full or empty, respectively).