

# Threads & Posix Threads (pthreads)

Z. Cihan TAYŞI



# Outline

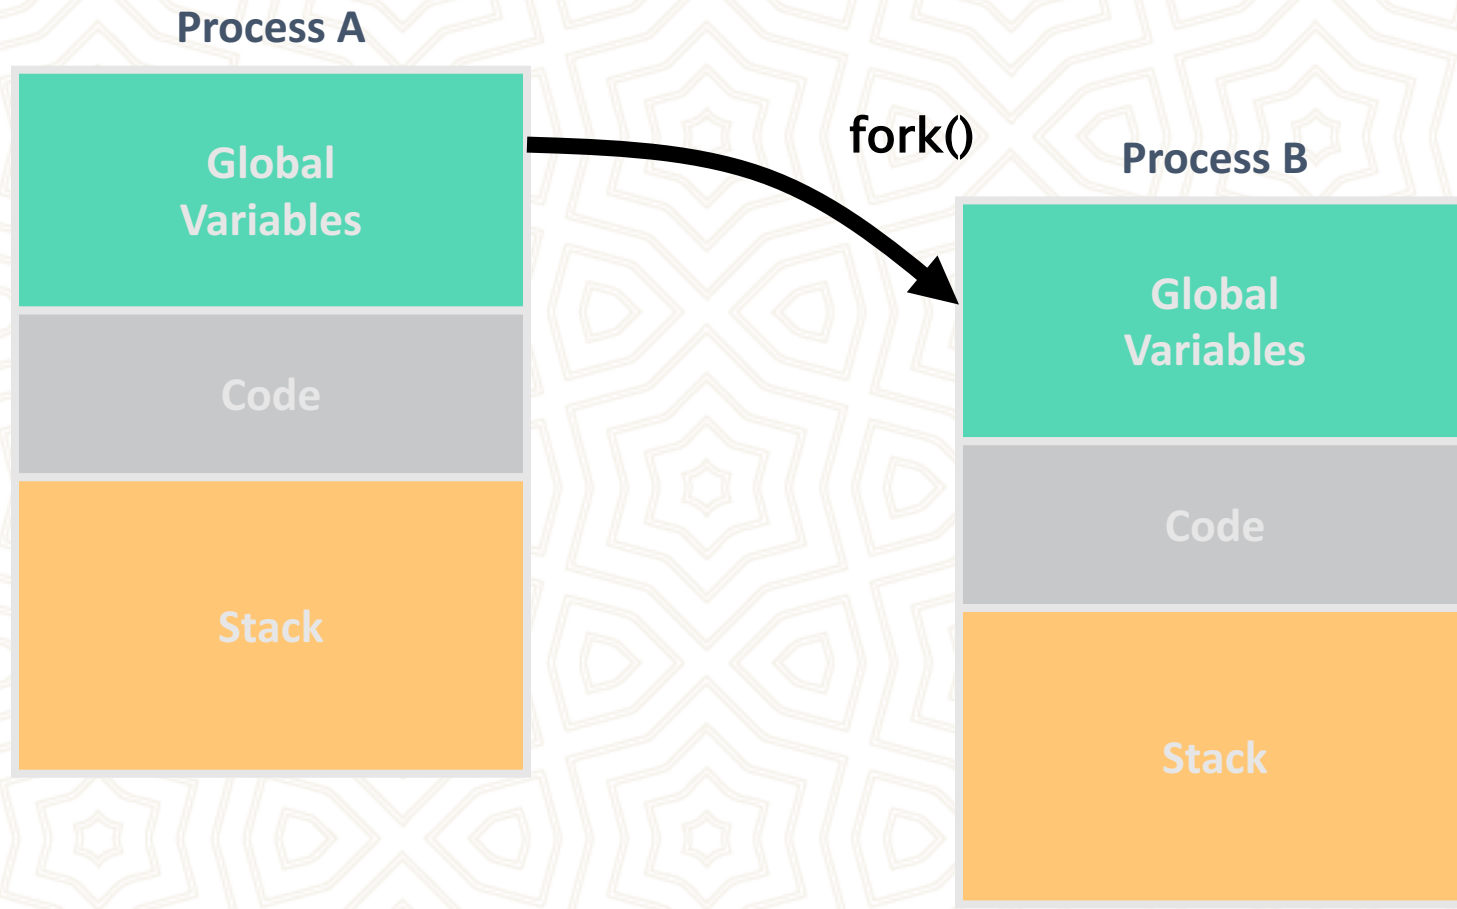
- Threads vs. Processes
- Multiple threads
- Thread-specific resources
- Posix threads
- Detached vs attached



# Threads vs. Processes

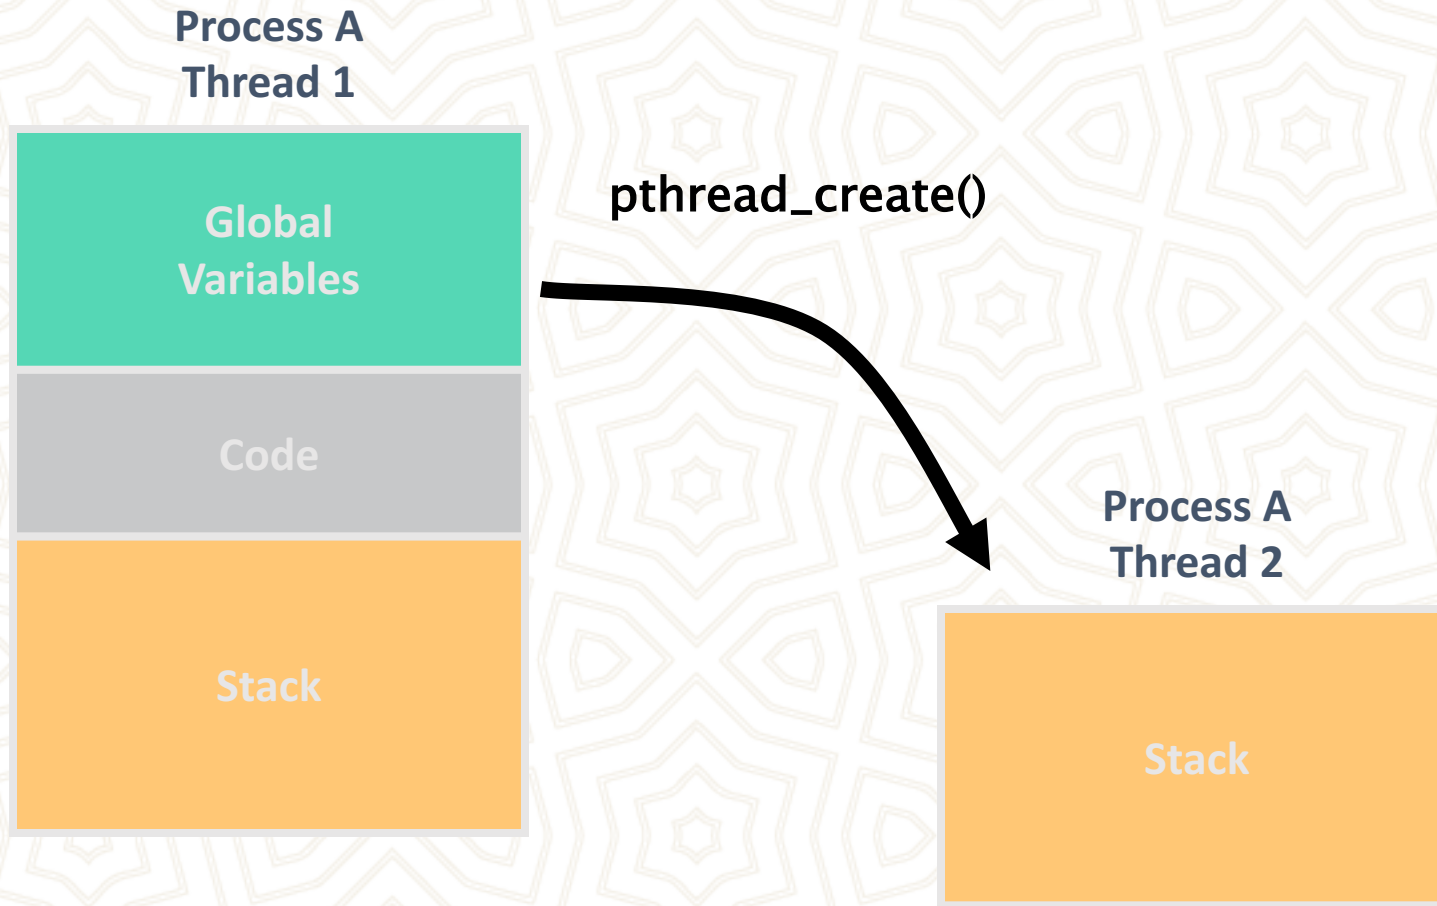
- Creation of a new process using fork is expensive
  - time & memory
- A thread does not require lots of memory or startup time.
  - sometimes called a lightweight process

# The fork() system call





# pthread\_create()



# Multiple Threads

- Each process can include many threads.
- All threads of a process share:
  - memory (program code and global data)
  - open file/socket descriptors
  - signal handlers and signal dispositions
  - working environment (current directory, user ID, etc.)



# Thread-specific Resources

- Each thread has its own
  - Thread ID
  - Stack, Registers, Program Counter
  - **errno** (if not - **errno** would be useless!)
- Threads within the same process can communicate using shared memory.
  - ***Must be done carefully***



# Posix Threads

- We will focus on Posix Threads - most widely supported threads programming API.
- you need to link with “-lpthread”
- On many systems this also forces the compiler to link in re-entrant libraries (instead of plain vanilla C libraries).





# Thread Creation

- `pthread_create(  
    pthread_t *tid,  
    const pthread_attr_t *attr,  
    void *(*func)(void *),  
    void *arg);`
- **func** is the function to be called.
  - when `func()` returns the thread is terminated.

# pthread\_create()

- The return value is 0 for OK.
  - **positive error number on error.**
- Does not set errno !!!
- Thread ID is returned in **tid**





# Thread IDs

- Each thread has a unique ID, a thread can find out it's ID by calling **pthread\_self()**.
- Thread IDs are of type `pthread_t` which is usually an unsigned int. When debugging, it's often useful to do something like this:
  - `printf("Thread %u:\n",pthread_self());`



# Thread Arguments

- When **func()** is called the value **arg** specified in the call to **pthread\_create()** is passed as a parameter.
- **func** can have **only 1** parameter, and it can't be larger than the size of a **void \***.



# Thread Arguments (cont.)

- Complex parameters can be passed by creating a structure and passing the address of the structure.
- The structure can't be a local variable (of the function calling **pthread\_create**)!!
  - threads have different stacks!



# A Simple pthread Example

pthread1.c

```
int main(int argc, char *argv[]) {

    pthread_t threads[NUM_THREADS];    // Thread identifiers
    int i, rc, *taskid[NUM_THREADS];   // Id numbers for each thread

    // Initialize the salutations for each thread
    messages[0] = "English: Hello World!";
    messages[1] = "French: Bonjour, le monde!";
    messages[2] = "Spanish: Hola al mundo";
    messages[3] = "Klingon: Nuq neH!";
    messages[4] = "German: Guten Tag, Welt!";

    for(i=0; i<NUM_THREADS; i++) {
        // Allocate an array for arguments to the threads
        taskid[i] = (int *) malloc(sizeof(int));
        *taskid[i] = i;
        // Create a thread with its argument in taskid[i]
        rc = pthread_create(&threads[i], NULL, PrintHello, (void *) taskid[i]);
        if (rc) { // Check for errors
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }
    pthread_exit(0);
}
```

```
void *PrintHello(void *threadid) {
    int *myarg;
    sleep(1); // Sleep for a second
    myarg = (int *) threadid; // Get own id from the argument
    printf("Thread %d: %s\n", *myarg, messages[*myarg]);
    pthread_exit(NULL);
}
```

```
lucid@ubuntu:~/Downloads/threads$ ./P1
Thread 2: Spanish: Hola al mundo
Thread 3: Klingon: Nuq neH!
Thread 4: German: Guten Tag, Welt!
Thread 1: French: Bonjour, le monde!
Thread 0: English: Hello World!
lucid@ubuntu:~/Downloads/threads$
```





# A Not So Simple pthread Example

pthread2.c

```
int main(int argc, char *argv[]) {
    pthread_t threads[NUM_THREADS];
    int rc, i, sum;

    sum=0;
    messages[0] = "English: Hello World!";
    messages[1] = "French: Bonjour, le monde!";
    messages[2] = "Spanish: Hola al mundo";
    messages[3] = "Klingon: Nuq neH!";
    messages[4] = "German: Guten Tag, Welt!";

    for(i=0; i<NUM_THREADS; i++) {
        // Initialize arguments to a thread
        sum = sum + i;
        thread_data_array[i].thread_id = i;
        thread_data_array[i].sum = sum;
        thread_data_array[i].message = messages[i];
        // Create a thread
        rc = pthread_create(&threads[i], NULL, PrintHello, &thread_data_array[i]);
        if (rc) {
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }
    pthread_exit(NULL);
}
```

```
void *PrintHello(void *threadarg) {
    int myid, sum;
    char *hello_msg;
    struct thread_data *my_data;

    sleep(1);
    my_data = (struct thread_data *) threadarg;
    myid = my_data->thread_id;
    sum = my_data->sum;
    hello_msg = my_data->message;
    printf("Thread %d: %s Sum=%d\n", myid, hello_msg, sum);
    pthread_exit(NULL);
}
```

```
lucid@ubuntu:~/Downloads/threads$ ./P2
Thread 3: Klingon: Nuq neH! Sum=6
Thread 4: German: Guten Tag, Welt! Sum=10
Thread 2: Spanish: Hola al mundo Sum=3
Thread 1: French: Bonjour, le monde! Sum=1
Thread 0: English: Hello World! Sum=0
lucid@ubuntu:~/Downloads/threads$
```



# Thread Lifespan

- ▶ Once a thread is created, it starts executing the function **func()** specified in the call to **pthread\_create()**.
- ▶ If **func()** returns, the thread is terminated.
- ▶ A thread can also be terminated by calling **pthread\_exit()**.
- ▶ If **main()** returns or any thread calls **exit()** all threads are terminated.



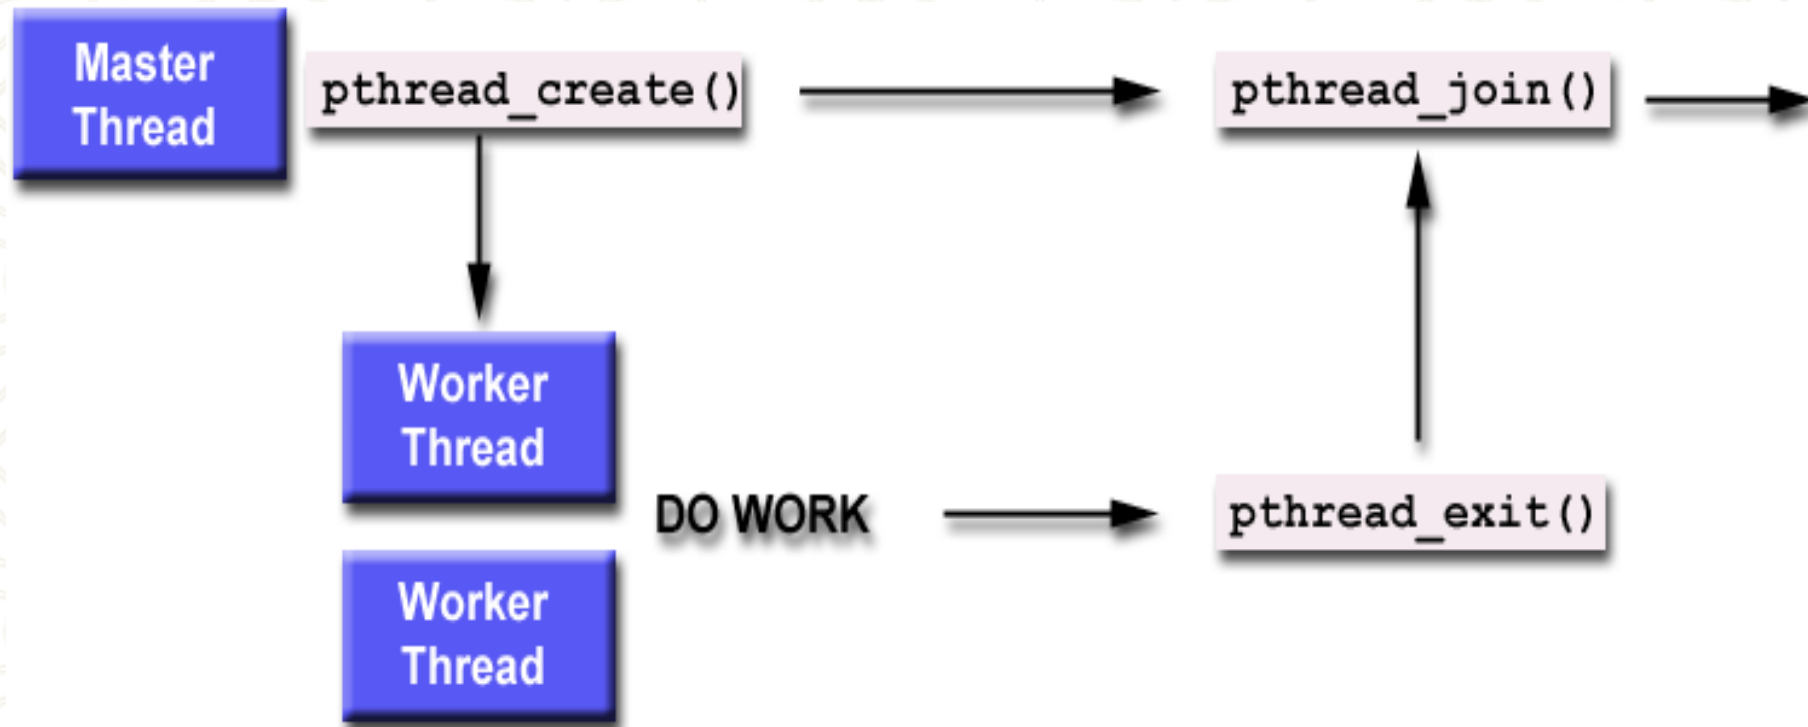


# Detached vs. Joinable

- Each thread can be either **joinable** or **detached**.
- **Joinable:** on thread termination the thread ID and exit status are saved by the OS.
- **Detached:** on termination all thread resources are released by the OS. A detached thread cannot be joined.



# Detached vs. Joinable (Contd.)





# Joinable Thread Example

```
// Initialize and set thread joinable attribute
pthread_attr_t attr;
pthread_attr_init(&attr);
pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_JOINABLE);

for(i=0; i<NUM_THREADS; i++) {
    rc = pthread_create(&thread[i], &attr, work, (void *)i);
    if (rc) {
        printf("ERROR; return code from pthread_create() is %d\n", rc);
        exit(-1);
    }
}

// Free attribute and join with worker threads
pthread_attr_destroy(&attr);
for(i=0; i<NUM_THREADS; i++) {
    rc = pthread_join(thread[i], &status); // Join with each thread
    if (rc) {
        printf("ERROR return code from pthread_join() is %d\n", rc);
        exit(-1);
    }
    res = *(double *)status; // Get the result from each thread
    total += res;             // Add it to total
    printf("Joined with thread %ld, result = %6.2f\n", i, res);
}
printf("\nProgram completed, total = %6.2f\n", total);
pthread_exit(NULL);
```

pthread3.c

```
void *work(void *t) {
    double *result;
    long int i, tid; // Pointers are 64 bit
    result = malloc(sizeof(double));
    *result = 0.0;
    tid = (long int)t; // Get the thread ID
    for (i=0; i<ITERATIONS; i++) {
        *result += (double)random()/RAND_MAX;
    }
    pthread_exit((void*) result); // Pass the result back
}
```

```
lucid@ubuntu:~/Downloads/threads$ ./P3
Joined with thread 0, result = 499951.58
Joined with thread 1, result = 500316.97
Joined with thread 2, result = 500138.36
Joined with thread 3, result = 500361.11
Joined with thread 4, result = 500214.25
Joined with thread 5, result = 499878.04
Joined with thread 6, result = 500303.43
Joined with thread 7, result = 499768.26
Joined with thread 8, result = 499445.38
Joined with thread 9, result = 499957.87
Program completed, total = 5000335.26
```



# Howto detach

```
#include <pthread.h>

pthread_t      tid;  // thread ID
pthread_attr_t attr; // thread attribute

// set thread detachstate attribute to DETACHED
pthread_attr_init(&attr);
pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_DETACHED);

// create the thread
pthread_create(&tid, &attr, start_routine, arg);
...
```





# Detached Thread Example

```
int main(int argc, char *argv[]) {
    pthread_t thread[NUM_THREADS];
    pthread_attr_t attr;
    int rc;
    long int t;

    // Initialize and set thread detached attribute
    pthread_attr_init(&attr);
    pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_DETACHED);

    for(t=0;t<NUM_THREADS;t++) {
        printf("Main: creating thread %ld\n", t);
        rc = pthread_create(&thread[t], &attr, BusyWork, (void *)t);
        if (rc) {
            printf("ERROR; return code from pthread_create() is %d\n", rc);
            exit(-1);
        }
    }

    // We're done with the attribute object, so we can destroy it
    pthread_attr_destroy(&attr);

    printf("Main: program completed. Exiting.\n");
    pthread_exit(NULL);
}
```

```
void *BusyWork(void *t) {
    long int i, tid;
    double maxval, result=0.0;
    tid = (long int)t;
    printf("Thread %ld starting...\n",tid);
    for (i=0; i<ITERATIONS; i++) {
        result += (double)random()/RAND_MAX;
    }
    printf("Thread %ld done. Result = %.2f\n",tid, result);
}
```

```
lucid@ubuntu:~/Downloads/threads$ ./P4
Main: creating thread 0
Main: creating thread 1
Main: creating thread 2
Main: creating thread 3
Main: program completed. Exiting.
Thread 2 starting...
Thread 1 starting...
Thread 3 starting...
Thread 0 starting...
Thread 3 done. Result = 499733.00
Thread 2 done. Result = 500148.38
Thread 0 done. Result = 500250.83
Thread 1 done. Result = 499750.66
lucid@ubuntu:~/Downloads/threads$
```

# Shared Global Variables

- Possible problems
  - Global variables
- Avoiding problems
- Synchronization Methods
  - Mutexes
  - Condition variables



# Possible problems

- **Sharing global variables is dangerous**
  - two threads may attempt to modify the same variable at the same time.
- Just because you don't see a problem when running your code doesn't mean it can't and won't happen!!!!

# Avoiding problems

- pthreads includes support for **Mutual Exclusion** primitives that can be used to protect against this problem.
- The general idea is to **lock** something before accessing global variables and to unlock as soon as you are done.
- **Shared socket descriptors** should be treated as **global variables!!!**





# Mutexes

- A global variable of type `pthread_mutex_t` is required:
- `pthread_mutex_t counter_mtx =  
PTHREAD_MUTEX_INITIALIZER;`
- Initialization to `PTHREAD_MUTEX_INITIALIZER` is required for a static variable!

# Lock & Unlock

- To lock use:
  - `pthread_mutex_lock(pthread_mutex_t &);`
- To unlock use:
  - `pthread_mutex_unlock(pthread_mutex_t &);`
- Both functions are blocking!



# Mutex Example

- A simple program to compute dot product of two vectors
  - a and b
- Main data is globally available to all threads
  - **dotdata** structure
    - dotdata.a
    - dotdata.b
  - Each threads works on a different part of it
- The sum is stored **dotdata.sum**
  - protected by a **mutex** !!



# Condition Variables

- **pthread**s support **condition variables**, which allow one thread to wait (sleep) for an event generated by any other thread.
- This allows us to avoid the **busy waiting** problem.
- `pthread_cond_t foo = PTHREAD_COND_INITIALIZER;`



# Condition Variables (cont.)

- A condition variable is always used with mutex.
- `pthread_cond_wait(pthread_cond_t *cptr,  
pthread_mutex_t *mptr);`
- `pthread_cond_signal(pthread_cond_t *cptr);`

*don't let the word signal confuse you -  
this has nothing to do with Unix signals*



# Condition Variable Example

- The main thread creates three threads.
  - Two of those threads increment a "count" variable,
  - The third thread watches the value of "count".
- When "**count**" reaches a predefined limit,
  - the waiting thread is signaled by one of the incrementing threads.



# Summary

- Threads are awesome, but dangerous. You have to pay attention to details or it's easy to end up with code that is incorrect (doesn't always work, or hangs in deadlock).
- Posix threads provides support for mutual exclusion, condition variables and thread-specific data.
- IHOP serves breakfast 24 hours a day!



# References

- <http://ube.ege.edu.tr/~cinsdiki/UBI511.html>
- man pages

