

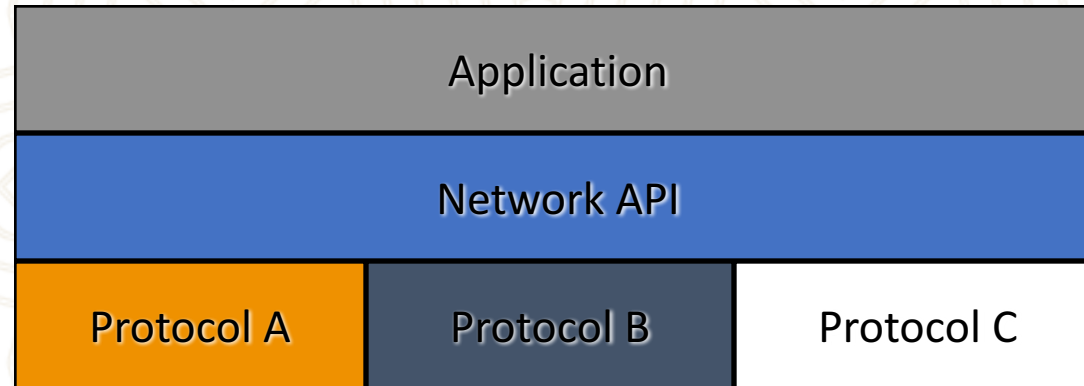
# Socket Programming

Z. Cihan TAYŞI



# Network Application Programming Interface (API)

- The services provided (often by the operating system) that provide the interface between application and protocol software.





# Network API

- Operating system provides Application Programming Interface (API) for network application
- **API is defined** by a set of function types, data structures, and constants
- **Desirable characteristics** of the network interface
  - Simple to use
  - Flexible
    - independent from any application
    - allows program to use all functionality of the network
  - Standardized
    - allows programmer to learn once, write anywhere
- Application Programming Interface for networks is called **socket**

# Sockets

- Sockets provide mechanisms to communicate between computers across a network
  - A socket is an abstract representation of a communication endpoint.
- There are different kind of sockets
  - DARPA Internet addresses (**Internet Sockets**)
  - Unix interprocess communication (Unix Sockets)
  - CCITT X.25 addresses
  - and many others
- **Berkeley sockets** is the most popular Internet Socket
  - runs on Linux, FreeBSD, OS X, Windows
  - fed by the popularity of TCP/IP



# Types of Internet Sockets

- Different types of sockets implement different communication types (stream vs. datagram)
- Type of socket: stream socket
  - connection-oriented
  - two way communication
  - reliable (error free), in order delivery
  - can use the Transmission Control Protocol (TCP)
  - e.g. telnet, ssh, http
- Type of socket: datagram socket
  - connectionless, does not maintain an open connection, each packet is independent
  - can use the User Datagram Protocol (UDP)
  - e.g. IP telephony
- Other types exist: similar to the one above

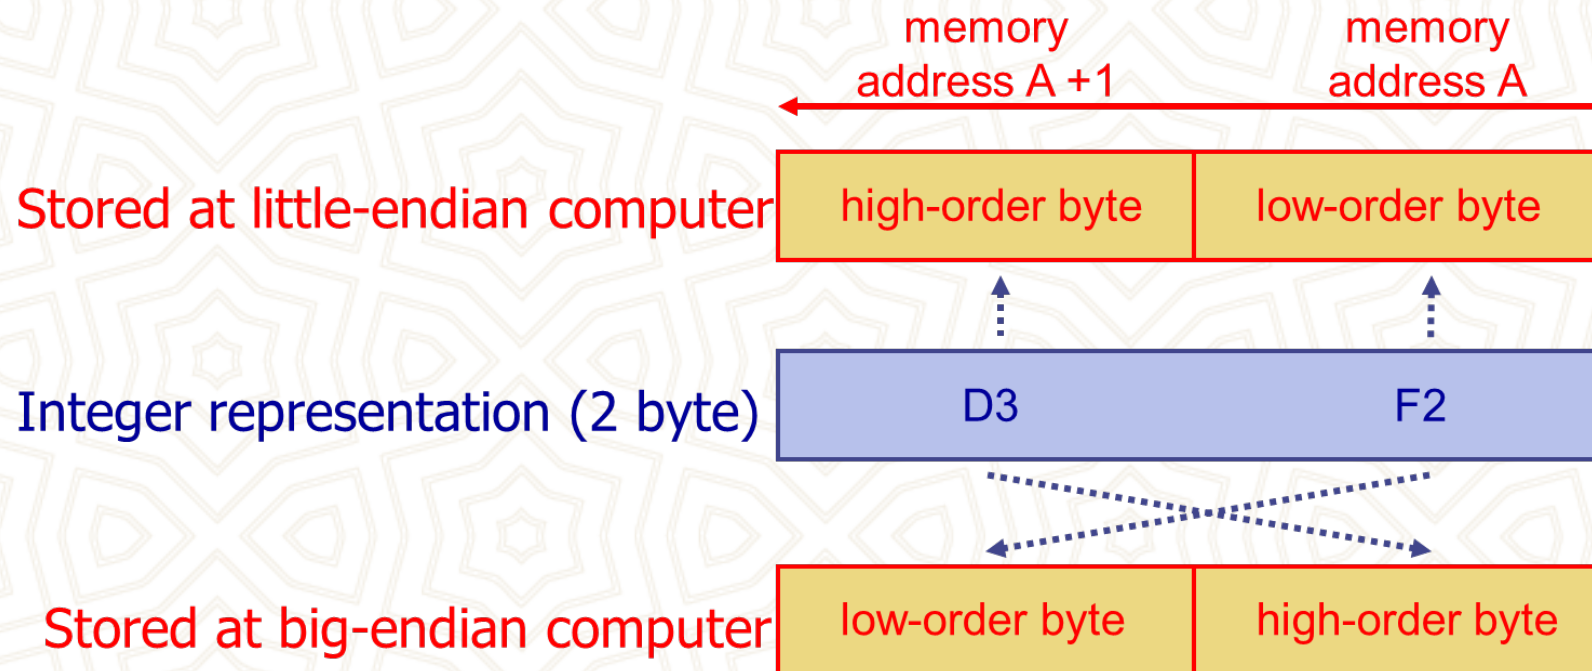
# Network Programming Tips

- Byte Ordering
- Naming
- Addressing



# Byte Ordering of Integers

- Different CPU architectures have different byte ordering



# Byte Ordering Problem

- **Question:** What would happen if two computers with different integer byte ordering communicate?
- **Answer:**
  - ◆ Nothing if they do not exchange integers!
  - ◆ But: If they exchange integers, they would get the wrong order of bytes, therefore, the wrong value!





# Byte Ordering Solution

- There are two solutions if computers with different byte ordering system want to communicate
  - They must **know the kind of architecture** of the sending computer (**bad solution**, it has not been implemented)
  - Introduction of a **network byte order**. The functions are:

```
uint16_t htons(uint16_t host16bitvalue)
uint32_t htonl(uint32_t host32bitvalue)
uint16_t ntohs(uint16_t net16bitvalue)
uint32_t ntohl(uint32_t net32bitvalue)
```

- Note: use for all integers (short and long), which are sent across the network
  - Including port numbers and IP addresses

# Name and Addressing

- **Host name**
  - identifies a single host (see Domain Name System slides)
  - **variable length string** (e.g. www.berkeley.edu)
  - is mapped to one or more IP addresses
- **IP Address**
  - written as dotted octets (e.g. 10.0.0.1)
  - **32 bits. Not a number! But often needs to be converted to a 32-bit to use.**
- **Port number**
  - identifies a process on a host
  - **16 bit number**

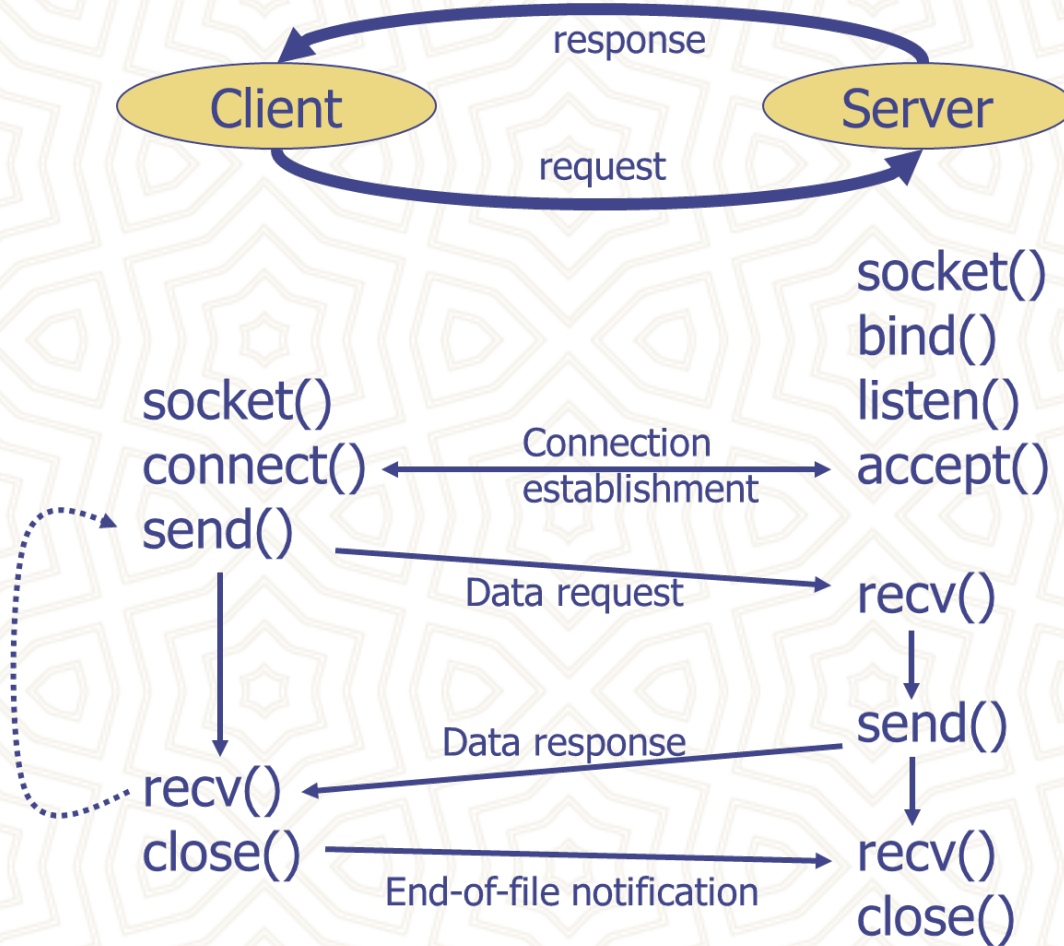


# Client-Server Architecture



- Client requests service from server
- Server responds with sending service or error message to client

# Simple Client-Server Example





# UDP Client – Server

## CLIENT

- Create stream socket
  - *socket()*
- While
  - *sendto()*
  - *recvfrom()*
- Close the Socket
  - *close()*

## SERVER

- Create stream socket
  - *socket()*
- Bind port to socket
  - *bind()*
- While
  - *recvfrom()*
  - *sendto()*
- Close the Socket
  - *close()*

# TCP Client – Server

## CLIENT

- Create stream socket
  - *socket()*
- Connect to server
  - *connect()*
- While still connected:
  - *send()*
  - *recv()*
- Close TCP connection and Socket
  - *close()*

## SERVER

- Create stream socket
  - *socket()*
- Bind port to socket
  - *bind()*
- Listen for new client
  - *listen()*
- While
  - *accept()*
  - *recv()*
  - *send()*
- Close TCP connection and Socket
  - *close()*



# Creating a Socket

```
int socket(int family,int type,int proto);
```

- family specifies the protocol family (**PF\_INET** for TCP/IP).
- type specifies the type of service (**SOCK\_STREAM**, **SOCK\_DGRAM**).
- protocol specifies the specific protocol
  - **AF\_INET, AF\_IPX, AF\_PACKET**

# Socket Descriptor Data Structure

**Descriptor Table**

0	•
1	•
2	•
3	•
4	•
	⋮


**Family: PF\_INET**  
**Service: SOCK\_STREAM**  
**Local IP: 111.22.3.4**  
**Remote IP: 123.45.6.78**  
**Local Port: 2249**  
**Remote Port: 3726**



# Assigning an address to a socket

- The **bind()** system call is used to assign an address to an existing socket.

```
int bind( int sockfd,  
          const struct sockaddr *myaddr,  
          const int addrlen);
```



- **bind** returns 0 if successful or -1 on error.

# Generic socket addresses

```
struct sockaddr {  
    sa_family_t    sa_family;  
    char           sa_data[14];  
};
```

- **sa\_family** specifies the address type.
- **sa\_data** specifies the address value.



# struct sockaddr\_in (IPv4)

```
struct sockaddr_in {  
    short int          sin_family;    // Address family  
    unsigned short int sin_port;      // Port number  
    struct in_addr      sin_addr;     // Internet address  
    unsigned char       sin_zero[8];  
};
```

```
struct in_addr {  
    unsigned long s_addr;    // 4 bytes  
};
```

- **Padding of `sin_zeros`:** `struct sockaddr_in` has same size as `struct sockaddr`

# Bind Example

```
int mysock, err;  
struct sockaddr_in myaddr;  
  
mysock = socket(PF_INET, SOCK_STREAM, 0);  
myaddr.sin_family = AF_INET;  
myaddr.sin_port = htons( portnum );  
myaddr.sin_addr = htonl( ipaddress);  
  
err=bind(mysock, (sockaddr *) &myaddr,  
        sizeof(myaddr));
```



# Sending UDP Datagrams

```
ssize_t sendto( int sockfd,  
                void *buff,  
                size_t nbytes,  
                int flags,  
                const struct sockaddr* to,  
                socklen_t addrlen);
```

**sockfd** is a UDP socket

**buff** is the address of the data (nbytes long)

**to** is the address of a sockaddr containing the destination address.

Return value is the number of bytes sent, or -1 on error.

# More **sendto()**

- The return value of **sendto** ( ) indicates how much data was accepted by the O.S. for sending as a datagram - not how much data made it to the destination.
- There is no error condition that indicates that the destination did not get the data!!!



# Receiving UDP Datagrams

```
ssize_t recvfrom( int sockfd,  
                  void *buff,  
                  size_t nbytes,  
                  int flags,  
                  struct sockaddr* from,  
                  socklen_t *fromaddrlen);
```

**sockfd** is a UDP socket

**buff** is the address of a buffer (**nbytes** long)

**from** is the address of a sockaddr.

Return value is the number of bytes received and put into buff, or -1 on error.

# More on **recvfrom()**

- If buff is not large enough, any extra data is lost forever...
- You can receive 0 bytes of data!
- The **sockaddr** at **from** is filled in with the address of the sender.
- You should set **fromaddrlen** before calling.
- If **from** and **fromaddrlen** are NULL we don't find out who sent the data.



# UDP Server

```
#define BUFLen 512
#define NPACK 10
#define PORT 9930

void diep(char *s) {
    perror(s);
    exit(1);
}

int main(void) {
    struct sockaddr_in si_me, si_other;
    int s, i, slen=sizeof(si_other);
    char buf[BUFLen];

    if ((s=socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP))==-1)
        diep("socket");

    memset((char *) &si_me, 0, sizeof(si_me));
    si_me.sin_family = AF_INET;
    si_me.sin_port = htons(PORT);
    si_me.sin_addr.s_addr = htonl(INADDR_ANY);
    if (bind(s, &si_me, sizeof(si_me))==-1)
        diep("bind");

    for (i=0; i<NPACK; i++) {
        if (recvfrom(s, buf, BUFLen, 0, (const struct sockaddr *) &si_other, &slen)==-1)
            diep("recvfrom()");
        printf("Received packet from %s:%d\nData: %s\n\n",
            inet_ntoa(si_other.sin_addr), ntohs(si_other.sin_port), buf);
    }

    close(s);
    return 0;
}
```

# UDP Client

```
#define BUFLen 512
#define NPACK 10
#define PORT 9930
#define SRV_IP "127.0.0.1"

void diep(char *s) {
    perror(s);
    exit(1);
}

int main(void) {
    struct sockaddr_in si_other;
    int s, i, slen=sizeof(si_other);
    char buf[BUFLen];

    if ((s=socket(AF_INET, SOCK_DGRAM, IPPROTO_UDP))==-1)
        diep("socket");

    memset((char *) &si_other, 0, sizeof(si_other));
    si_other.sin_family = AF_INET;
    si_other.sin_port = htons(PORT);
    if (inet_aton(SRV_IP, &si_other.sin_addr)==0) {
        fprintf(stderr, "inet_aton() failed\n");
        exit(1);
    }

    for (i=0; i<NPACK; i++) {
        printf("Sending packet %d\n", i);
        sprintf(buf, "This is packet %d\n", i);
        if (sendto(s, buf, BUFLen, 0, (const struct sockaddr *) &si_other, slen)==-1)
            diep("sendto()");
    }

    close(s);
    return 0;
}
```



# Wait for Connections **listen()**

**int listen(int sockfd, int backlog);**

- Puts socket in a listening state, willing to handle incoming TCP connection request.
- Backlog: number of TCP connections that can be queued at the socket.

# Accept Connections **accept()**

```
int accept(int sockfd, struct sockaddr *addr, socklen_t *addrlen);
```

- The **accept()** system call is used with connection-based socket types (**SOCK\_STREAM, SOCK\_SEQPACKET**).
- It extracts the first connection request on the queue of pending connections for the listening socket, *sockfd*, creates a new connected socket, and returns a new file descriptor referring to that socket.
- The newly created socket is not in the listening state. The original socket *sockfd* is unaffected by this call.



# Sending Packets - **send()**

```
int send_packets(char *buffer, int buffer_len) {  
    sent_bytes = send(chat_sock, buffer, buffer_len, 0);  
    if (sent_bytes < 0) {  
        perror ("send");  
    }  
    return 0;  
}
```

- Needs socket descriptor,
- Buffer containing the message, and
- Length of the message
- Can also use **write()**

# Receiving packets **recv()**

```
ssize_t recv(int sockfd, void *buf, size_t len, int flags);
```

Wait for a maximum of length octets of data on the **SOCK\_STREAM** socket sockfd ? write data to buffer



# A Simple TCP Client – Server

- A Hello World example
  - Server listens at a predefined port
  - Accepts incoming connections
  - Sends Hello World string to clients
- Extra features
  - Address resolution

# TCP Client

```
#include <arpa/inet.h>
#include <netinet/in.h>
#include <stdio.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <unistd.h>
#include <string.h>
#include <stdlib.h>
#include <netinet/in.h>
#include <sys/socket.h>
#include <netdb.h>

#define PORT 3490          // the port client will be connecting to
#define MAXDATASIZE 100    // max number of bytes we can get
                           // at once

int main(int argc, char *argv[]) {
    int sockfd, numbytes;
    char buf[MAXDATASIZE];
    struct hostent *he;
    struct sockaddr_in their_addr; // server's address information

    if (argc != 2) {
        fprintf(stderr, "usage: client hostname\n");
        exit(1);
    }

    if ((he=gethostbyname(argv[1])) == NULL) { // get the host info
        perror("gethostbyname");
        exit(1);
    }

    if ((sockfd = socket(AF_INET, SOCK_STREAM, 0)) == -1) {
        perror("socket");
        exit(1);
    }
}
```



# TCP Client

```
their_addr.sin_family = AF_INET;           // host byte order
their_addr.sin_port = htons(PORT);         // short, network byte order
their_addr.sin_addr = *((struct in_addr *)he->h_addr); // already network byte order
memset(&(their_addr.sin_zero), '\0', 8);    // zero the rest of the struct

if (connect(sockfd, (struct sockaddr *)&their_addr, sizeof(struct sockaddr)) == -1){
    perror("connect");
    exit(1);
}

if ((numbytes=recv(sockfd, buf, MAXDATASIZE-1, 0)) == -1) {
    perror("recv");
    exit(1);
}

buf[numbytes] = '\0';
printf("Received: %s", buf);
close(sockfd);
return 0;
}
```

# TCP Server

```
#include <arpa/inet.h>
#include <netinet/in.h>
#include <stdio.h>
#include <sys/types.h>
#include <sys/socket.h>
#include <unistd.h>
#include <string.h>
#include <stdlib.h>

#define MYPORT 3490 // the port users will be connecting to
#define BACKLOG 10 // how many pending connections queue will hold

int main(void) {
    |
    int sockfd, new_fd; // listen on sockfd, new connection on new_fd
    struct sockaddr_in my_addr; // my address information
    struct sockaddr_in their_addr; // connector's address information
    socklen_t sin_size;

    if ((sockfd = socket(AF_INET, SOCK_STREAM, 0)) == -1) { // perror("socket");
        exit(1);
    }

    my_addr.sin_family = AF_INET; // host byte order
    my_addr.sin_port = htons(MYPORT); // short, network byte order
    my_addr.sin_addr.s_addr = INADDR_ANY; // auto. filled with local IP
    memset(&(my_addr.sin_zero), '\0', 8); // zero the rest of the struct
}
```



# TCP Server

```
if (bind(sockfd, (struct sockaddr *)&my_addr, sizeof(struct sockaddr)) == -1) {
    perror("bind");
    exit(1);
}

if (listen(sockfd, BACKLOG) == -1) {
    perror("listen");
    exit(1);
}

while(1) {    // main accept() loop
    sin_size = sizeof(struct sockaddr_in);
    if ((new_fd = accept(sockfd, (struct sockaddr *)&their_addr, &sin_size)) == -1) {
        perror("accept");
        continue;
    }

    printf("server: got connection from %s\n", inet_ntoa(their_addr.sin_addr));

    if (send(new_fd, "Hello, world!\n", 14, 0) == -1)
        perror("send");

    close(new_fd);
}
return 0;
```

# More things to know


- Howto send broadcast packages ?
  - UDP
  - TCP
- Howto implement non-blocking socket ?
  - UDP
  - TCP
- Howto implement **concurrent** servers



# Assignment II

- Connection oriented Client-Server application
  - Server broadcasts its information
    - IP Address and port number
    - When receives a packet from the client;
      - it takes system time
      - append this information to the packet
      - sends it back.
  - Clients listens
    - When the server information is available, it connects to the server
    - When connected, it sends a packet with time information
    - When receives the packet back
      - gets time information from the packet
      - Calculate delay for both packets
      - Calculate the turn around time

# Assignment II





# References

- Man pages
- Jörn Altmann's Slides
- *UNIX Network Programming, Volume 1, Second Edition: Networking APIs: Sockets and XTI*