

Real-Time Systems

Week 1 - Concept

What is a real-time system ?

- ▶ A real-time system is a system that must complete its tasks and provide its service in a time sensitive manner.
- ▶ It typically must monitor, respond to, or control an external environment.
- ▶ It is connected to its environment through sensors, actuators, and other input-output devices.
- ▶ Must meet timing constraints imposed by external environment.

Examples of Real-time systems

- ▶ Vehicle systems
 - ▶ ABS breaking, power steering, dynamic traction control
- ▶ Traffic control
 - ▶ Traffic lights, air traffic control
- ▶ Process control
 - ▶ Power plants, chemical processes, consumer applications
- ▶ Communication systems
 - ▶ Voice and video transmission, data transmission
- ▶ Medical systems
- ▶ Automated manufacturing

Jobs & Tasks

- ▶ As we wish to characterize a wide range of real-time systems, we want to discuss the work done in general terms.
- ▶ Every unit of work scheduled and executed on system is referred to as **a job**.
- ▶ We refer to a set of related jobs that together perform a system function, as **a task**.
- ▶ So, we might have a task T_i , composed of k jobs
 - ▶ $J_{i,1}, J_{i,2}, \dots, J_{i,k}$

Release time

- ▶ The instant a job is available to be executed, we refer to this as its **release time**.
- ▶ A job can't be scheduled and executed before its release time.
- ▶ After its release time, must also wait for any data or control dependencies before it can be scheduled and executed.

Release time

- ▶ Example: Say a system monitors and controls several furnace.
 - ▶ After it starts executing at $t = 0$, it reads each temperature sensor and stores the value in memory every 100ms.
 - ▶ It also computes the control law for each furnace every 100ms, with the first computation occurring at $t = 20$ ms.
- ▶ We can capture the fact control laws are computed periodically in terms of release time.
 - ▶ i.e. Jobs J_0, J_1, \dots, J_k have release times of $20 + k \times 100$ ms.

Deadlines

- ▶ A job's **deadline** is point in time that its execution must be completed by.
- ▶ If a job has no deadline, we say its deadline is at infinity.
- ▶ We refer to a job's **response time** as the length of time between its release time, and moment it completes.
- ▶ A job's maximum allowed response time is called its **relative deadline**.
- ▶ A job's **deadline is thus equal to its release time plus relative deadline**.

Hard and soft timing constraints

- ▶ A **hard real-time** timing constraint is a constraint that causes system failure, if violated.
 - ▶ controlling the motion of an elevator
 - ▶ stopping a train before a collision, that a track switch position before train reaches switch
- ▶ A **soft real-time** timing constraint is a constraint that allows the system to operate effectively if a few deadlines are missed.
 - ▶ voice transmission data during a telephone call (a few packets dropped probably won't be missed)
 - ▶ electronic games (if a frame is slightly delayed now and then, not a big deal.)

Reference model

- ▶ A real-time system is characterized by three things:
 - ▶ **Workload model:** describes the system's applications.
 - ▶ **Resource model:** describes the system's resources.
 - ▶ This includes processors as well as passive resources such as memory, mutexes, etc.
 - ▶ **Scheduling algorithms**
 - ▶ Describes how the system makes use of its resources.

Workload models

- ▶ Periodic tasks
 - ▶ a deterministic workload model.
- ▶ Aperiodic tasks
 - ▶ have soft or no deadlines.
- ▶ Sporadic tasks
 - ▶ have hard deadlines.

Periodic tasks - I

- ▶ A **periodic task** is a computation or data transfer that occurs on a regular or semiregular basis in order to provide an ongoing function of the system.
- ▶ Each periodic task T_i is composed of constituent jobs
 - ▶ $J_{i,1}, J_{i,2}, \dots$ where $J_{i,k}$ is the k th job in T_i
- ▶ The release time, or phase, of task T_i is the release time ($r_{i,1}$) of $J_{i,1}$, and is labelled (Φ_i)
- ▶ The period of task T_i is (p_i). This means the jobs of T_i get released every p_i time units.

Periodic tasks - II

- ▶ We use (D_i) to refer to the relative deadline of task T_i , and (d_i) to refer to the absolute deadline.
- ▶ We often assume that $D_i = \Phi_i$,
 - ▶ a job from T_i has time till the next job is released to complete.
- ▶ Sometimes we have $D_i < \Phi_i$ in order to reduce jitter (variation in the completion times).

Periodic tasks - II

- ▶ The **execution time** (e_i) of task T_i is the maximum time it takes for any one of its jobs to complete.
- ▶ The **utilization** of task T_i , $u_i = e_i/\Phi_i$, is the fraction of time the task keeps a processor busy if it executes for its maximum time for each job.
- ▶ Total utilization of processor is
 - ▶ $u = \sum_{i=1}^n (e_i/\Phi_i)$

Aperiodic and Sporadic tasks


- ▶ Jobs in a non-periodic task arrive (are released) according to some probability distribution $A(x)$.
 - ▶ This is the probability that the time between releases is x .
- ▶ The execution time of the task is given by some probability distribution $B(x)$.
- ▶ The time between successive job arrivals can be arbitrarily small.
- ▶ **Aperiodic tasks** have soft or no deadlines. Our goal is to optimize response time of these tasks, but not at expense of hard real-time tasks.
- ▶ **Sporadic tasks** have hard deadlines.

Precedence Constraints

- ▶ Jobs that can execute in any order are said to be independent.
- ▶ Jobs have precedence constraints when they are constrained to execute in a given order.
 - ▶ We use the relation $J_i < J_k$ to indicate that job J_k may not start until job J_i has completed.
 - ▶ We say that J_i is a predecessor of J_k , and that J_k is a successor of J_i .

Precedence Constraints

- ▶ A job is ready to execute after its release time and when all of its predecessors have completed.
- ▶ Directed graph $G = (J, <)$ is called a precedence graph and shows the predecessor relationships.



Preemptivity of jobs

- ▶ Execution of jobs can occur in interleaved fashion.
- ▶ Execution of less urgent job could be suspended in order to give processor to a more urgent job.
- ▶ When job completes, processor is returned to less urgent job to continue execution.
- ▶ When we interrupt job execution like this, it's called **preemption**.
- ▶ We say a job is **preemptable** if it can be suspended at any time, allow other jobs to execute, and then be resumed from suspension point.

Scheduler

- ▶ A **scheduler** implements a set of scheduling algorithms and resource access-control protocols used to decide when individual jobs are to be executed and what resources they should be allocated.
- ▶ In particular, a scheduler assigns jobs to processors for execution.
- ▶ We assume that a scheduler will always produce a **valid schedule**.

Scheduler

- ▶ A valid schedule satisfies the following conditions:
 - ▶ 1. Every process is assigned at most one job at a time.
 - ▶ 2. Every job is assigned to at most one processor at a time.
 - ▶ 3. No job is scheduled before its release time.
 - ▶ 4. The amount of processor time assigned to every job is equal to actual execution of job or its maximum execution time.
 - ▶ 5. All precedence and resource usage constraints are satisfied.

Terminology for Schedules

- ▶ A **schedule** is an assignment of all jobs in system on the processors, which are available.
- ▶ A **feasible schedule** is one where all jobs complete by their deadlines.
- ▶ A group of jobs are **schedulable** according to an algorithm, if algorithm always produces a feasible schedule.
- ▶ A scheduling algorithm is **optimal** if the scheduling algorithm always produces a feasible schedule when one exists.

Performance measures

- ▶ **Tardiness**
 - ▶ Difference between a job's completion time and its deadline, zero if deadline met.
 - ▶ Interested in maximum and average tardiness.
- ▶ **Lateness**
 - ▶ Difference between a job's completion time and its deadline, negative if early.
- ▶ **Response time**
 - ▶ The length of time between a job's release time, and moment it completes.

Performance measures

- ▶ **Miss rate**
 - ▶ Percentage of jobs executed but complete too late.
- ▶ **Loss rate**
 - ▶ Percentage of jobs discarded.
- ▶ **Invalid rate**
 - ▶ Percentage of jobs that either complete too late or are discarded.

Next week ...

- ▶ Scheduling algorithms
 - ▶ clock-driven approaches
 - ▶ priority driven approaches
 - ▶ preemption
 - ▶ online vs offline
 - ▶ ...