

Chapter 6: Information Retrieval and Web Search

Dr. Mehmet S. Aktaş

Acknowledgement: Thanks to Dr. Bing Liu for teaching materials.

Introduction

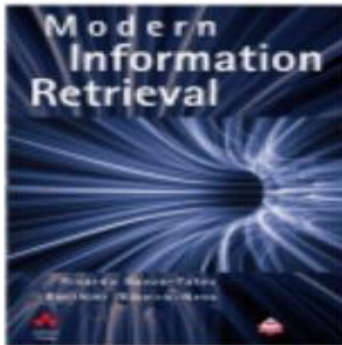
- Text mining refers to data mining using text documents as data.
- Most text mining tasks use **Information Retrieval** (IR) methods to pre-process text documents.
- These methods are quite different from traditional data pre-processing methods used for relational tables.
- Web search also has its root in IR.

Information Retrieval (IR)

- Conceptually, IR is the study of finding needed information. I.e., IR helps users find information that matches their information needs.
 - Expressed as queries
- Historically, IR is about document retrieval, emphasizing document as the basic unit.
 - Finding documents relevant to user queries
- Technically, IR studies the acquisition, organization, storage, retrieval, and distribution of information.



What is Information Retrieval (IR)?



IR: Part of computer science which studies the **retrieval of information (not data)** from a collection of **written documents**. The retrieved documents aim at satisfying a **user information need** usually expressed in **natural language**.

- Documents, unstructured, text, large
- Information need
- Store, search, find
- The World Wide Web?
- Relational databases?

DIKW

- **Data:** Raw web pages
 - **Information:** Result of query
 - **Knowledge:** Result of processing query result by user
 - **Wisdom:** Synthesis of many such actions by a set of users
-
- One possible classification of steps in process



Information Retrieval vs. Databases

Information retrieval	Data retrieval
Retrieve all objects relevant to some information need	Retrieve all objects satisfying some clearly defined conditions
Find all documents about the topic "semantic web"!	SELECT id FROM document WHERE title LIKE '%semantic web%'
Result list	Well-defined result set



```
[selke@tddb ~]$ db2 "SELECT id FROM document WHERE title LIKE '%semantic web%' FETCH FIRST 3 ROWS ONLY"
```

```
ID
-----
          45489
          9635899
          98556

3 record(s) selected.
```



Web Search

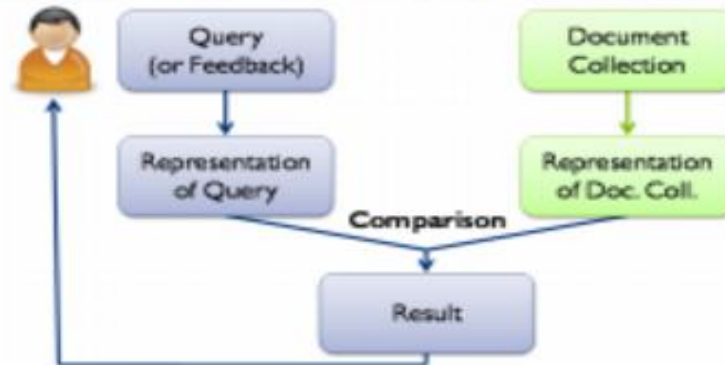
- Very similar to information retrieval
- Main differences:
 - **Links** between Web pages can be exploited
 - **Collecting**, storing, and **updating** documents is more difficult
 - Usually, the **number of users** is very large
 - **Spam** is a problem



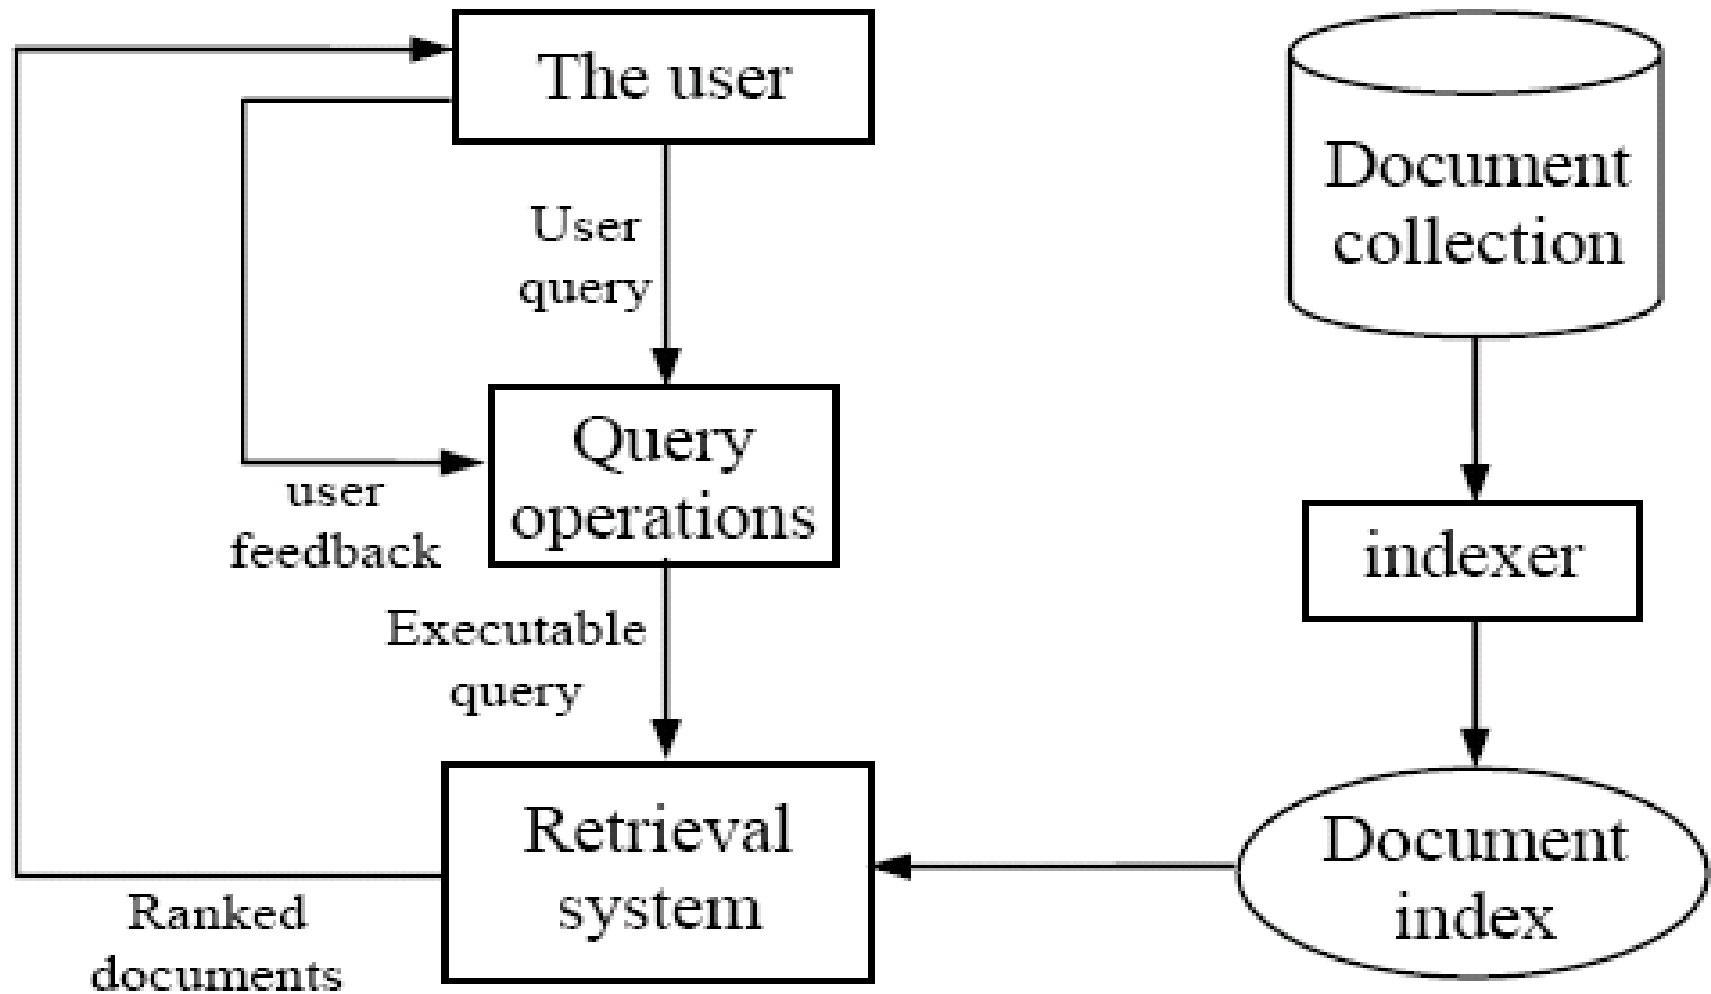


IR Models

- Any IR system is based on an **IR model**
- The model defines ...
 - ... a **query language**,
 - ... an internal **representation of queries**,
 - ... an internal **representation of documents**,
 - ... a **ranking function** which associates a real number with each query–document pair.
- Optional: A mechanism for **relevance feedback**



IR architecture



IR queries

- Keyword queries
- Boolean queries (using AND, OR, NOT)
- Phrase queries
- Proximity queries
- Full document queries
- Natural language questions

Information retrieval models

- An IR model governs how a document and a query are represented and how the relevance of a document to a user query is defined.
- Main models:
 - Boolean model
 - Vector space model
 - Statistical language model
 - etc

Boolean model

- Each document or query is treated as a **“bag” of words** or **terms**. Word sequence is not considered.
- Given a collection of documents D , let $V = \{t_1, t_2, \dots, t_{|V|}\}$ be the set of distinctive words/terms in the collection. V is called the **vocabulary**.
- A weight $w_{ij} > 0$ is associated with each term t_i of a document $\mathbf{d}_j \in D$. For a term that does not appear in document \mathbf{d}_j , $w_{ij} = 0$.

$$\mathbf{d}_j = (w_{1j}, w_{2j}, \dots, w_{|V|j}),$$

Boolean model (contd)

- Query terms are combined logically using the Boolean operators **AND**, **OR**, and **NOT**.
 - E.g., *((data AND mining) AND (NOT text))*
- Retrieval
 - Given a Boolean query, the system retrieves every document that makes the query logically true.
 - Called **exact match**.
- The retrieval results are usually quite poor because term frequency is not considered.

Vector space model

- Documents are also treated as a “bag” of words or terms.
- Each document is represented as a vector.
- However, the term weights are no longer 0 or 1. Each term weight is computed based on some variations of **TF** or **TF-IDF** scheme.
- **Term Frequency (TF) Scheme:** The weight of a term t_i in document \mathbf{d}_j is the number of times that t_i appears in \mathbf{d}_j , denoted by f_{ij} . Normalization may also be applied.

TF-IDF term weighting scheme

- The most well known weighting scheme

- TF: still **term frequency**
- IDF: **inverse document frequency**.

N : total number of docs

df_i : the number of docs that t_i appears.

- The final TF-IDF term weight is:

$$tf_{ij} = \frac{f_{ij}}{\max\{f_{1j}, f_{2j}, \dots, f_{|V|j}\}}$$

$$idf_i = \log \frac{N}{df_i}$$

$$w_{ij} = tf_{ij} \times idf_i.$$

Retrieval in vector space model

- Query \mathbf{q} is represented in the same way or slightly differently.
- **Relevance of \mathbf{d}_j to \mathbf{q}** : Compare the similarity of query \mathbf{q} and document \mathbf{d}_j .
- Cosine similarity (the cosine of the angle between the two vectors)

$$\text{cosine}(\mathbf{d}_j, \mathbf{q}) = \frac{\langle \mathbf{d}_j \bullet \mathbf{q} \rangle}{\|\mathbf{d}_j\| \times \|\mathbf{q}\|} = \frac{\sum_{i=1}^{|V|} w_{ij} \times w_{iq}}{\sqrt{\sum_{i=1}^{|V|} w_{ij}^2} \times \sqrt{\sum_{i=1}^{|V|} w_{iq}^2}}$$

- Cosine is also commonly used in text clustering

An Example

- A document space is defined by three terms:
 - hardware, software, users
 - the vocabulary
- A set of documents are defined as:
 - $A1=(1, 0, 0), \quad A2=(0, 1, 0), \quad A3=(0, 0, 1)$
 - $A4=(1, 1, 0), \quad A5=(1, 0, 1), \quad A6=(0, 1, 1)$
 - $A7=(1, 1, 1) \quad A8=(1, 0, 1). \quad A9=(0, 1, 1)$
- If the Query is “hardware and software”
- what documents should be retrieved?

An Example (cont.)

■ In Boolean query matching:

- document A4, A7 will be retrieved (“AND”)
- retrieved: A1, A2, A4, A5, A6, A7, A8, A9 (“OR”)

■ In similarity matching (cosine):

- $q=(1, 1, 0)$
- $S(q, A1)=0.71$, $S(q, A2)=0.71$, $S(q, A3)=0$
- $S(q, A4)=1$, $S(q, A5)=0.5$, $S(q, A6)=0.5$
- $S(q, A7)=0.82$, $S(q, A8)=0.5$, $S(q, A9)=0.5$
- Document retrieved set (with ranking)=
 - $\{A4, A7, A1, A2, A5, A6, A8, A9\}$

Okapi relevance method

- Another way to assess the degree of relevance is to directly compute a relevance score for each document to the query.
- The **Okapi** method and its variations are popular techniques in this setting.

The Okapi relevance score of a document d_j for a query q is:

$$okapi(d_j, q) = \sum_{t_i \in q, d_j} \ln \frac{N - df_i + 0.5}{df_i + 0.5} \times \frac{(k_1 + 1)f_{ij}}{k_1(1 - b + b \frac{dl_j}{avdl}) + f_{ij}} \times \frac{(k_2 + 1)f_{iq}}{k_2 + f_{iq}},$$

where k_1 (between 1.0-2.0), b (usually 0.75) and k_2 (between 1-1000)

Relevance feedback

- Relevance feedback is one of the techniques for improving retrieval effectiveness. The steps:
 - the user first identifies some relevant (D_r) and irrelevant documents (D_{ir}) in the initial list of retrieved documents
 - the system expands the query \mathbf{q} by extracting some additional terms from the sample relevant and irrelevant documents to produce \mathbf{q}_e
 - Perform a second round of retrieval.
- **Rocchio method** (α , β and γ are parameters)

$$\mathbf{q}_e = \alpha \mathbf{q} + \frac{\beta}{|D_r|} \sum_{\mathbf{d}_r \in D_r} \mathbf{d}_r - \frac{\gamma}{|D_{ir}|} \sum_{\mathbf{d}_{ir} \in D_{ir}} \mathbf{d}_{ir}$$

Text pre-processing

- Word (term) extraction: easy
- Stopwords removal
- Stemming
- Frequency counts and computing TF-IDF term weights.

Stopwords removal

- Many of the most frequently used words in English are useless in IR and text mining – these words are called *stop words*.
 - ❑ the, of, and, to,
 - ❑ Typically about 400 to 500 such words
 - ❑ For an application, an additional domain specific stopwords list may be constructed
- Why do we need to remove stopwords?
 - ❑ Reduce indexing (or data) file size
 - stopwords accounts 20-30% of total word counts.
 - ❑ Improve efficiency and effectiveness
 - stopwords are not useful for searching or text mining
 - they may also confuse the retrieval system.

Stemming

- Techniques used to find out the root/stem of a word. E.g.,

- ❑ user
- ❑ users
- ❑ used
- ❑ using

engineering
engineered
engineer

- stem: use

engineer

Usefulness:

- improving effectiveness of IR and text mining
 - ❑ matching similar words
 - ❑ Mainly improve recall
- reducing indexing size
 - ❑ combining words with same roots may reduce indexing size as much as 40-50%.

Basic stemming methods

Using a set of rules. E.g.,

■ remove ending

- ❑ if a word ends with a consonant other than s, followed by an s, then delete s.
- ❑ if a word ends in es, drop the s.
- ❑ if a word ends in ing, delete the ing unless the remaining word consists only of one letter or of th.
- ❑ If a word ends with ed, preceded by a consonant, delete the ed unless this leaves only a single letter.
- ❑

■ transform words

- ❑ if a word ends with “ies” but not “eies” or “aies” then “ies --> y.”

Frequency counts + TF-IDF

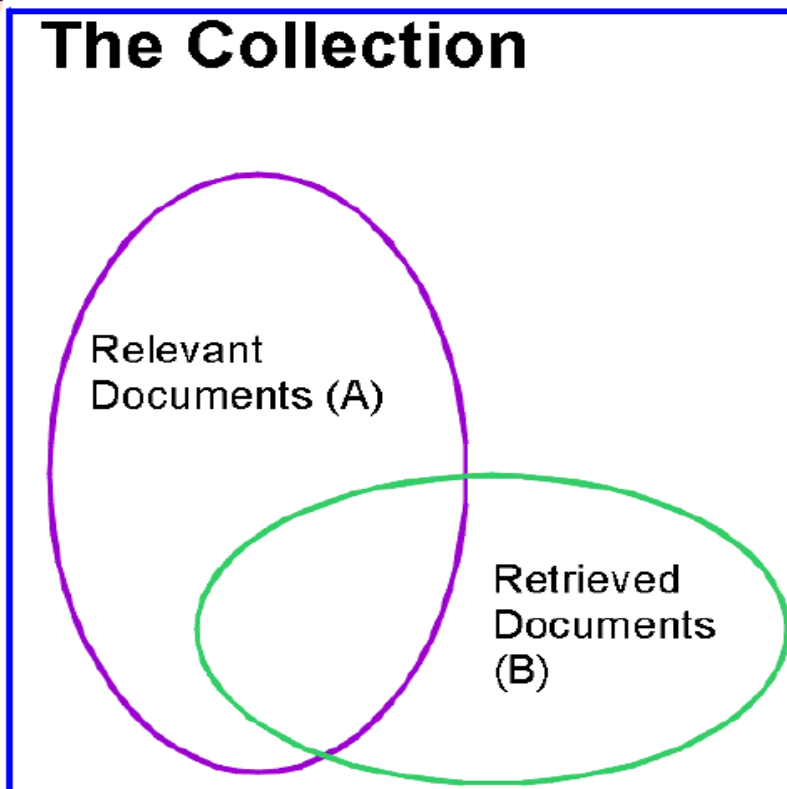
- Counts the number of times a word occurred in a document.
 - Using occurrence frequencies to indicate relative importance of a word in a document.
 - if a word appears often in a document, the document likely “deals with” subjects related to the word.
- Counts the number of documents in the collection that contains each word
- TF-IDF can be computed.

Evaluation: Precision and Recall

- Given a query:
 - Are all retrieved documents relevant?
 - Have all the relevant documents been retrieved?
- Measures for system performance:
 - The first question is about the **precision** of the search
 - The second is about the completeness (**recall**) of the search.

"Classic" Information Retrieval

The Collection



- Given a query, the system retrieves a set B of documents
- Every retrieved document is either relevant or irrelevant to the query

Quality metrics:

- Recall : $(A \cap B) / A$
- Precision : $(A \cap B) / B$



Recall and Precision on the Web

- Relevance of document to queries is not binary – there are many shades of gray
- Broad-topic queries:
 - abundance problem
 - Precision is the dominating factor: users mostly satisfied with a few good results (a few *authoritative* pages)
- Narrow-topic queries:
 - Find a needle in an enormous haystack
 - Recall demands engines cover significant portions of the Web
- Common measure: precision@10
- Nowadays larger emphasis on *diversity*
 - Positive recall for many aspects of the query

Precision-recall curve

Example 2: Following Example 1, we obtain the interpolated precisions at all 11 recall levels in the table of Fig. 6.4. The precision-recall curve is shown on the right.

i	$p(r_i)$	r_i
0	100%	0%
1	100%	10%
2	100%	20%
3	100%	30%
4	80%	40%
5	80%	50%
6	71%	60%
7	70%	70%
8	70%	80%
9	62%	90%
10	62%	100%

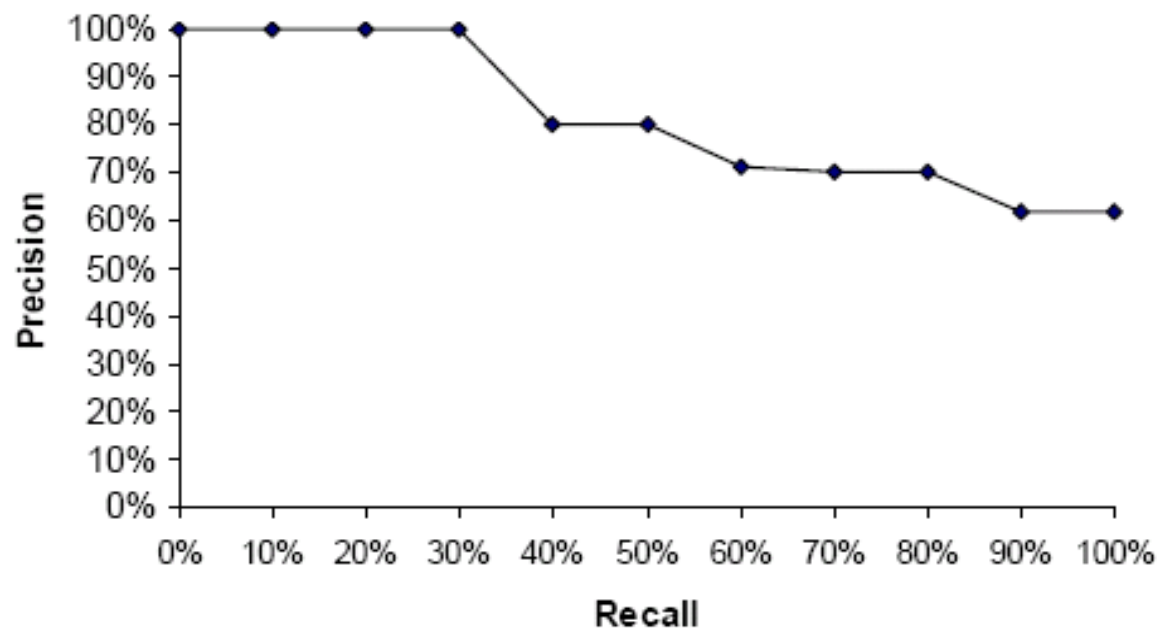


Fig. 6.4. The precision-recall curve

Compare different retrieval algorithms

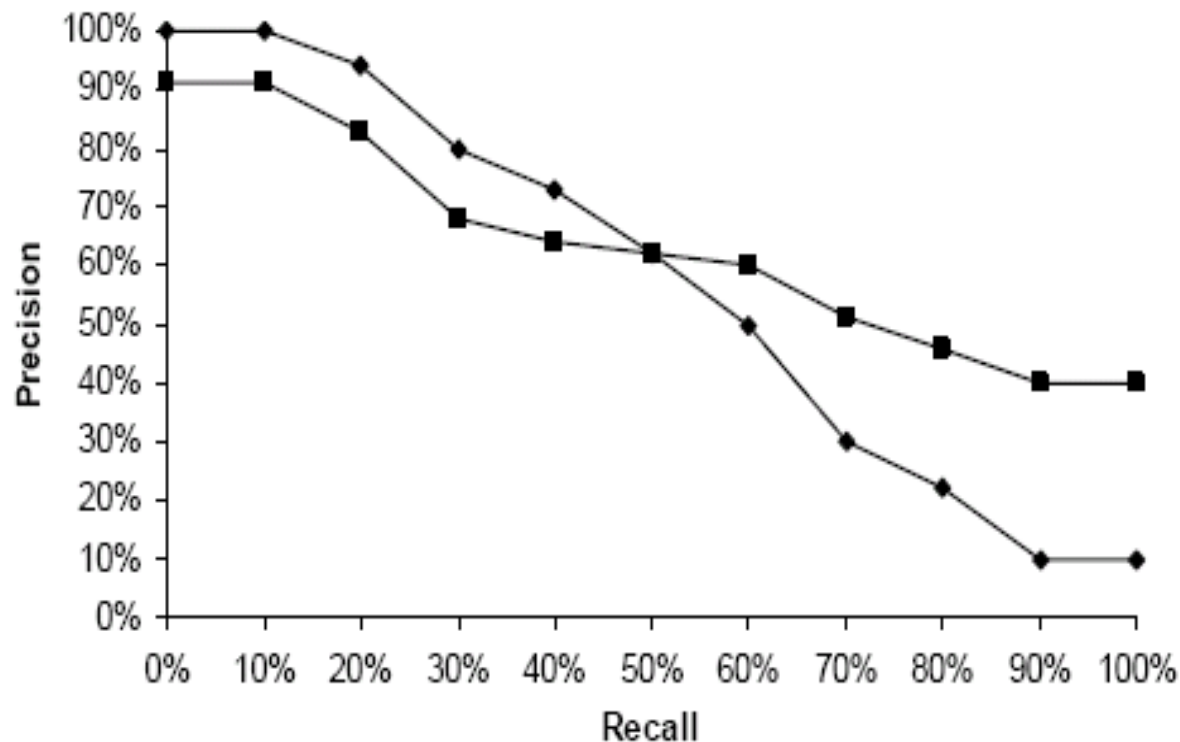


Fig. 6.5. Comparison of two retrieval algorithms based on their precision-recall curves

Compare with multiple queries

- Compute the average precision at each recall level.

$$\bar{p}(r_i) = \frac{1}{|Q|} \sum_{j=1}^{|Q|} p_j(r_i), \quad (22)$$

where Q is the set of all queries and $p_j(r_i)$ is the precision of query j at the recall level r_i . Using the average precision at each recall level, we can also draw a precision-recall curve.

- Draw precision recall curves

Rank precision

- Compute the precision values at some selected rank positions.
- Mainly used in Web search evaluation.
- For a Web search engine, we can compute precisions for the top 5, 10, 15, 20, 25 and 30 returned pages
 - as the user seldom looks at more than 30 pages.
- Recall is not very meaningful in Web search.
 - Why?

Web Search as a huge IR system

- A Web crawler (robot) crawls the Web to collect all the pages.
- Servers establish a huge inverted indexing database and other indexing databases
- At query (search) time, search engines conduct different types of vector query matching.

Inverted index

- The inverted index of a document collection is basically a data structure that
 - attaches each distinctive term with a list of all documents that contains the term.
- Thus, in retrieval, it takes constant time to
 - find the documents that contains a query term.
 - multiple query terms are also easy handle as we will see soon.

An example

Example 3: We have three documents of id_1 , id_2 , and id_3 :

id_1 : Web mining is useful.

1 2 3 4

id_2 : Usage mining applications.

1 2 3

id_3 : Web structure mining studies the Web hyperlink structure.

1 2 3 4 5 6 7 8

Applications: id_2

Hyperlink: id_3

Mining: id_1, id_2, id_3

Structure: id_3

Studies: id_3

Usage: id_2

Useful: id_1

Web: id_1, id_3

(A)

Applications: $\langle id_2, 1, [3] \rangle$

Hyperlink: $\langle id_3, 1, [7] \rangle$

Mining: $\langle id_1, 1, [2] \rangle, \langle id_2, 1, [2] \rangle, \langle id_3, 1, [3] \rangle$

Structure: $\langle id_3, 2, [2, 8] \rangle$

Studies: $\langle id_3, 1, [4] \rangle$

Usage: $\langle id_2, 1, [1] \rangle$

Useful: $\langle id_1, 1, [4] \rangle$

Web: $\langle id_1, 1, [1] \rangle, \langle id_3, 2, [1, 6] \rangle$

(B)

Fig. 6.7. Two inverted indices: a simple version and a more complex version

Search using inverted index

Given a query q , search has the following steps:

- Step 1 (**vocabulary search**): find each term/word in q in the inverted index.
- Step 2 (**results merging**): Merge results to find documents that contain all or some of the words/terms in q .
- Step 3 (**Rank score computation**): To rank the resulting documents/pages, using,
 - content-based ranking
 - link-based ranking

Different search engines

- The real differences among different search engines are
 - their index weighting schemes
 - Including location of terms, e.g., title, body, emphasized words, etc.
 - their query processing methods (e.g., query classification, expansion, etc)
 - **their ranking algorithms**
 - Few of these are published by any of the search engine companies. They are tightly guarded secrets.

Summary

- We only give a **VERY** brief introduction to IR. There are a large number of other topics, e.g.,
 - Statistical language model
 - Latent semantic indexing (LSI and SVD).
 - (read an IR book or take an IR course)
- Many other interesting topics are not covered, e.g.,
 - Web search
 - Index compression
 - Ranking: combining contents and hyperlinks
 - Web page pre-processing
 - Combining multiple rankings and meta search
 - Web spamming
- Want to know more? Read the textbook