

A First Parallel Program

Chapter 4

Primality Testing

A simple computation that will take a long time.

Whether a number x is prime: Decide whether a number x is prime using the trial division algorithm.

Trial Division Algorithm: The function tries to divide x by 2 and by every odd number p from 3 up to the square root of x .

If any remainder is 0, then p is a factor of x and x is not prime; otherwise x is prime.

Primality Testing Function

```
public static boolean isPrime(long x) {
    if (x % 2 == 0) {
        return false;
    }

    long p = 3;
    long xsqrt = (long) Math.ceil(Math.sqrt(x));
    while (p <= xsqrt) {
        if (x % p == 0) {
            return false;
        }
        p += 2;
    }
    return true;
}
```

Sequential Program

First, call this method sequentially for many numbers and keep the running times.

Measuring Time in Java

To measure times, we use Java's `System.currentTimeMillis()` method, which returns the wall clock time in milliseconds (msec) since 1970.

No printing while measuring time

We record each instant in a variable, and postpone printing the results, so as to disturb the timing as little as possible while the program is running. It can take several msec to call `println()`, and we don't want to include that time in our measurements.

```
/*
 * @author maktas
 */
public class PrimeTestSeq {

    /**
     * only test odd divisors.
     * even numbers can not be primes
     * start with 3 and test with every odd number: 3, 5, 7, 9,
     * go up to the square root of the tested number
     * the numbers larger than the square root can not divide be sole
divisors
     */
    private static boolean isPrime(long x) {
        if (x % 2 == 0) {
            return false;
        }

        long p = 3;
        long xsqrt = (long) Math.ceil(Math.sqrt(x));
        while (p <= xsqrt) {
            if (x % p == 0) {
                return false;
            }
            p += 2;
        }
        return true;
    }

    public static void main(String[] args) {
//        testOneNumber();
        testManyNumbers();
    }

    public static void testOneNumber() {
        long m = 1289237867378231L;
        System.out.println("started testing: " + m);

        long t1 = System.currentTimeMillis();
        boolean positive = isPrime(m);
        long t2 = System.currentTimeMillis();
        long duration = t2 - t1;

        if (positive) {
            System.out.println(m + " is prime");
        } else {
            System.out.println(m + " is not prime");
        }
        System.out.println("the time it takes: " + duration + " ms");
    }
}
```

```

public static void testManyNumbers() {
    long m[] = {10000000000000037L, 10000000000000091L,
100000000000000159L, 100000000000000187L};
    System.out.println("started testing: ");
    boolean positives[] = new boolean[m.length];
    long startTimes[] = new long[m.length];
    long endTimes[] = new long[m.length];

    long start = System.currentTimeMillis();

    for (int i = 0; i < m.length; i++) {
        startTimes[i] = System.currentTimeMillis();
        positives[i] = isPrime(m[i]);
        endTimes[i] = System.currentTimeMillis();
    }

    for (int i = 0; i < m.length; i++) {
        if (positives[i]) {
            System.out.println(m[i] + " is prime");
        } else {
            System.out.println(m[i] + " is not prime");
        }
        System.out.println(i + " start: " + (startTimes[i]-start) );
        System.out.println(i + " end: " + (endTimes[i]-start) );
    }
}
}

```

Iterative Program Running Times




Figure 4.1 Program1Seq execution timeline, SMP parallel computer

Parallel Program

Parallel Java library

Parallel java library jar file needs to be downloaded and included into the project.

The file to be downloaded: **pj20120620.jar**

Address: <http://www.cs.rit.edu/~ark/pj.shtml#download>

Add the jar file to: Project | Properties | Libraries

ParallelTeam object

A ParallelTeam object is constructed. This object runs the threads in parallel. We provide the number of threads to be run to this object.

ParallelRegion: run method

ParallelRegion object has the code that will run in parallel in the run method.

Anonymous class

ParallelRegion object constructed from an anonymous class.

getThreadIndex() method

the index of the calling thread is retrieved by ParallelRegion's getThreadIndex()

static variables

the variables that need to be accessed by all threads are defined as static. They are defined outside of **main** method.

Arrays

Since each thread needs to access its part of the variables. We define each variable as arrays. Each thread accesses the array element by its index.

startTimes, endTimes, numbers, variables are arrays.

Array sizes are equal to the number of threads.

```
/*
 * @author maktas
 */
import edu.rit.pj.ParallelRegion;
import edu.rit.pj.ParallelTeam;

public class PrimeTestParallel {

    /**
     * only test odd divisors.
     * even numbers can not be primes
     * start with 3 and test with every odd number: 3, 5, 7, 9,
     * go up to the square root of the tested number
     * the numbers larger than the square root can not divide be sole
     divisors
     */
}
```

```

private static boolean isPrime(long x) {
    if (x % 2 == 0) {
        return false;
    }

    long p = 3;
    long xsqrt = (long) Math.ceil(Math.sqrt(x));
    while (p <= xsqrt) {
        if (x % p == 0) {
            return false;
        }
        p += 2;
    }
    return true;
}

static long start;
static long startTimes[];
static long endTimes[];
static long numbers[] = {10000000000000037L, 10000000000000091L};
// static long numbers[] = {10000000000000037L, 10000000000000091L,
10000000000000159L, 10000000000000187L};

public static void main(String[] args) throws Exception {

    start = System.currentTimeMillis();
    startTimes = new long[numbers.length];
    endTimes = new long[numbers.length];

    ParallelTeam team = new ParallelTeam(numbers.length);
    team.execute(new ParallelRegion() {

        public void run() {
            int i = getThreadIndex();
            startTimes[i] = System.currentTimeMillis();
            isPrime(numbers[i]);
            endTimes[i] = System.currentTimeMillis();
        }
    });

    for (int i = 0; i < numbers.length; i++) {
        System.out.println(i + " start: " + (startTimes[i]-start) );
        System.out.println(i + " end:   " + (endTimes[i]-start) );
    }
}
}

```

How Parallel Program Works

Main program thread: The main program begins with one thread, the “main thread,” executing the main() method.

Parallel Team Threads: When the main thread creates the parallel team object, the parallel team object creates additional hidden threads; the constructor argument specifies the number of threads.

Team work: These form a “team” of threads for executing code in parallel. When the main thread calls the parallel region’s execute() method, the main thread suspends execution and the parallel team threads take over.

All team threads call ParallelRegion’s run method simultaneously: Each thread retrieves a value for x, and each thread calls isPrime().

Concurrent execution: Thus, the isPrime() subroutine calls happen at the same time, and each subroutine call is performed by a different thread with a different argument.

Back to main thread: When all the subroutine calls have finished executing, the main thread resumes executing statements after the parallel region and prints the timing measurements.

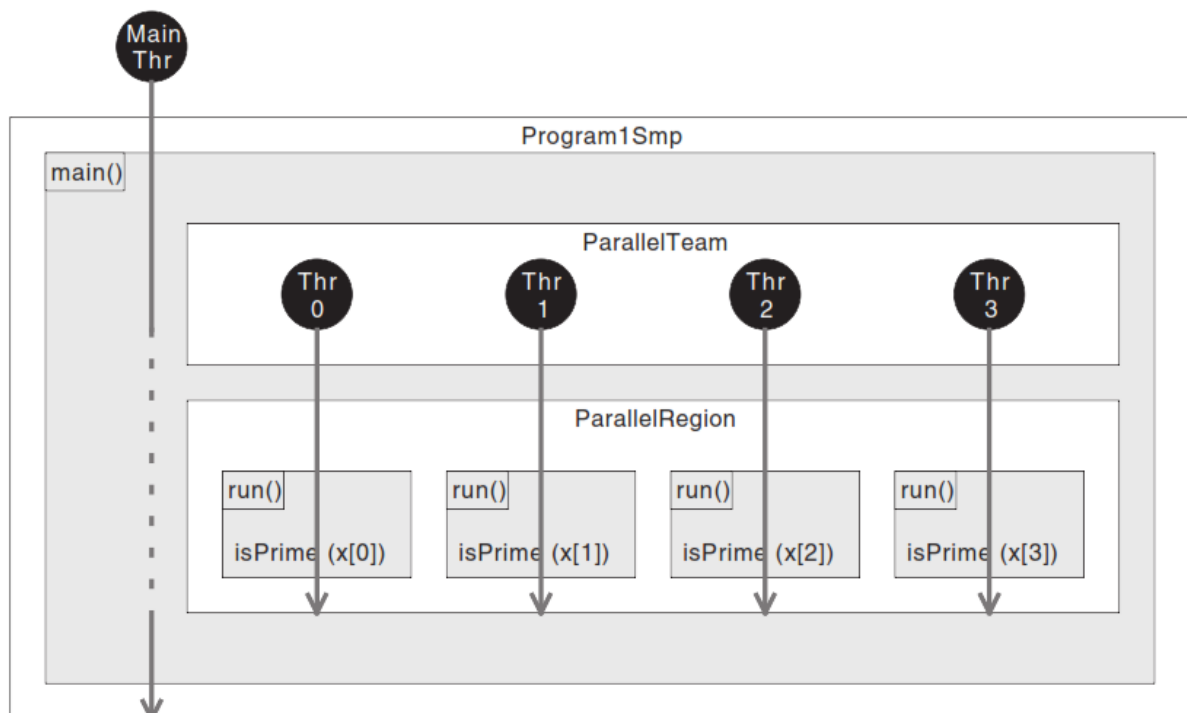


Figure 4.2 Program1Smp operation

Thread to Processor Assignments

When running such a thread-based program on an SMP parallel computer, the Java Virtual Machine (JVM) and the operating system are responsible for scheduling each thread to execute on a different processor.

No Loop, Multiple Threads

The parallel program illustrates a central theme of parallel program design: Repetition does not necessarily imply sequencing (looping).

The sequential program used a loop to get n repetitions (4 repetitions) of a subroutine call (isPrime).

However, for this program there is no need to do the repetitions in sequence in a loop. Each method call is assigned to a thread.

So a loop is not the only way to do a repeated calculation.

Running Timelines

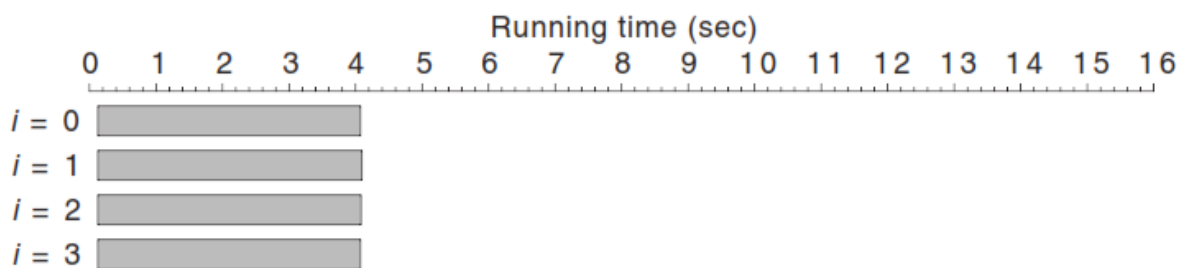


Figure 4.3 Program1Smp execution timeline, SMP parallel computer

Each thread runs in parallel. There needs to be at least 4 cores.

The speedup from the book:

Sequential running time/parallel running time = $15342/4098 = 3.744$

Almost linear increase.

Startup Cost for parallel programs

With Program1Smp, the first subroutine call didn't begin until 125 msec after the program started.

During this time, the program was occupied in creating the parallel team and parallel region objects, starting up the parallel team threads, and executing the parallel region's run() method—work that the sequential program didn't have to do.

Parallelism is not free

This illustrates another central theme of parallel program design: Parallelism is not free. The benefit of speedup or sizeup comes with a price of extra overhead that is not needed in a sequential program.

The name of the game is to minimize this extra overhead.

Running Parallel Programs on a Single CPU Machine

A parallel program can easily run on a single CPU machine.

In a single CPU machine, parallel version may outperform the sequential version.

This is surprising since parallel program also have threading overhead.

The reason for this is that:

JVM Just-In-Time compiler converts the isPrime function into machine code faster in parallel version. Since all threads are calling the same function.

In sequential version, it takes longer time for JVM to understand that isPrime function is a hotspot and its code needs to be converted into the machine code.

Parallel Programming without Parallel Java Library

We can write the same program using Java threads and not using Parallel java library.

In this case, we explicitly construct and control the threads. We start them and monitor when they are done.

Parallel Java makes many things much simpler. Particularly the thread synchronization and communication.

```
/**
 * Primality test with threads and without using Parallel Java library
 *
 * @author maktas
 */
public class PrimeTestThreads extends Thread{

    int index;

    public PrimeTestThreads(int index){
        this.index = index;
    }
}
```



```

public void run() {
    startTimes[index] = System.currentTimeMillis();
    PrimeTestParallel02.isPrime(numbers[index]);
    endTimes[index] = System.currentTimeMillis();
}

/**
 * only test odd divisors.
 * even numbers can not be primes
 * start with 3 and test with every odd number: 3, 5, 7, 9,
 * go up to the square root of the tested number
 * the numbers larger than the square root can not divide be sole
divisors
 */
public static boolean isPrime(long x) {
    if (x % 2 == 0) {
        return false;
    }

    long p = 3;
    long xsqrt = (long) Math.ceil(Math.sqrt(x));
    while (p <= xsqrt) {
        if (x % p == 0) {
            return false;
        }
        p += 2;
    }
    return true;
}

static long start;
static long startTimes[];
static long endTimes[];
static long numbers[] = {10000000000000037L, 10000000000000091L};
// static long numbers[] = {10000000000000037L, 10000000000000091L,
10000000000000159L, 10000000000000187L};

public static void main(String[] args) throws Exception {

    start = System.currentTimeMillis();
    startTimes = new long[numbers.length];
    endTimes = new long[numbers.length];

    PrimeTestThreads threadler[] = new
PrimeTestThreads[numbers.length];
    for (int i = 0; i < threadler.length; i++) {
        threadler[i] = new PrimeTestThreads(i);
    }

    for (int i = 0; i < threadler.length; i++) {
        threadler[i].start();
    }

    // wait until all threads are done
    boolean notDone = true;
    while(notDone){

        notDone = false;
        for (int i = 0; i < threadler.length; i++) {
            if(threadler[i].isAlive())

```

```
        notDone = true;
    }
    try{
        Thread.sleep(100);
    }catch(Exception e){

    }
}

for (int i = 0; i < numbers.length; i++) {
    System.out.println(i + " start: " + (startTimes[i] - start));
    System.out.println(i + " end:   " + (endTimes[i] - start));
}
}
```